

BADM 371 Intro to Data Analytics

BADM 371

2022-12-14

Contents

1	Introduction	7
2	Syllabus	9
2.1	Course Objectives and Learning Outcomes	9
2.2	Text and Resources	10
2.3	Performance Evaluation (Grading)	10
2.4	Class Participation:	11
2.5	Phones	11
2.6	Attendance	11
2.7	Honor Code and Plagiarism:	12
2.8	First-Generation Version:	12
2.9	Continuing Advocate Version	12
2.10	Justice, Equity, Diversity, and Inclusion Syllabus Statements . . .	12
2.11	Accessible Education	13
2.12	Bias Reporting	13
2.13	Counseling and Health	13
2.14	Presby First+ Program	13
2.15	Title IX Syllabus Statement	14
3	Schedule	15
	Spring 2023	15

4	R basics and workflows	17
4.1	Basics of working with R at the command line and RStudio goodies	17
4.2	In Rstudio - where we will live	17
4.3	17
4.4	Workspace and working directory	21
4.5	Exercises	23
5	Objects and Arithmetic	25
5.1	Introduction	25
5.2	Basic Functions	26
5.3	Statistics and Summaries	28
5.4	Exercises	30
6	Summaries and Subscripting	31
6.1	Introduction	31
6.2	Exercises (Summaries and Subscripting)	32
7	Matrices	33
7.1	CBind and RBind	33
7.2	Matrix Function	34
7.3	Exercises	36
8	Data Frames	39
8.1	Excercises	40
9	Preloaded data and mtcars	43
9.1	Practicing with mtcars data set	43
9.2	Excerises for you:	46
10	More simple data wrangling	49
10.1	a nice, fun little matrix for you	49
10.2	More fun (this class is really awesome isn't it?)	49

<i>CONTENTS</i>	5
11 GDH Ice Cream	53
11.1 Problem Introduction	53
11.2 Assignment	53
12 Exam 1 Practice	55
13 Quick Linear Regression	59
13.1 Quick Linear regression using Loblolly	59
14 Linear Regression with mtcars	63
14.1 Excercises for you	63
14.2 More Excercises for you	64
15 Deeper Linear Regression	65
15.1 Interpretation of the Model	65
15.2 Applying the Model to Make Predictions	67
15.3 Review Questions	69
16 Linear Regression Practice	71
16.1 Your homework is to watch these videos which are posted under the linear regression header on Brightspace and then do the following homework:	71
16.2 Problems	71
16.3 Multivariate Regression	72
17 Filters and packages	73
17.1 R packages	73
17.2 You can install the latest released version from CRAN with: . . .	74
17.3 In RStudio	74
17.4 Check this out	74
18 dplyr	75
18.1 Introduction	75
18.2 A Neat Resource	75
18.3 Single table verbs	75

18.4 The pipe	76
18.5 Loading <code>dplyr</code>	76
18.6 <code>starwars</code> examples	76
19 Practice	79
19.1 <code>starwars</code> Exercises	79
19.2 <code>iris</code> Exercises	79
19.3 More practice	80
19.4 Yet more practice	81
20 <code>dplyr</code> and <code>nycflights13</code>	83
20.1 Choosing columns: <code>select</code> , <code>rename</code>	83
20.2 Choosing rows: <code>filter</code> , <code>between</code> , <code>slice</code> , <code>sample_n</code> , <code>top_n</code> , <code>distinct</code> .	86
20.3 Exercises	90
20.4 Adding new variables: <code>mutate</code> , <code>transmute</code> , <code>add_rownames</code>	90
20.5 Grouping and counting: <code>summarise</code> , <code>tally</code> , <code>count</code> , <code>group_size</code> , <code>n_groups</code> , <code>ungroup</code>	92
20.6 Creating data frames: <code>data_frame</code>	96
21 Joining tables	97
21.1 Joining (merging) tables: <code>left_join</code> , <code>right_join</code> , <code>inner_join</code> , <code>full_join</code> , <code>semi_join</code> , <code>anti_join</code>	97
21.2 Viewing more output: <code>print</code> , <code>View</code>	99
22 More DPLYR Exercises	103
23 NFL Data	105
23.1 Additional practice	105
24 Palmer Penguins Review	107
24.1 Use the <code>penguins</code> dataset from the <code>palmerpenguins</code> package for the following:	107
25 Diamonds Practice	109
25.1 Use the <code>diamonds</code> data set found in the <code>ggplot2</code> package to an- swer these questions. You may want to review the data before starting using the help options.	109

Chapter 1

Introduction

This is a book written which contains all the materials and lessons for **BADM 371 Intro to Data Analytics**. If you miss a day, want to review something we covered in class, or for any reason want to look for a worksheet, this is where you can go.

Each chapter contains a different topic we will cover during the semester. Some larger topics are split into two chapters to make accessing the materials a little more intuitive.

This book is updated automatically with any changes made to the documents during the semester, so if at any point you are told there was a change in the assignment, you can come here to get the updated version.

Also, this book has benefited greatly from lots of free, readily available resources posted on the web and we leverage these extensively. I would encourage you to review these resources in your analytics journey. Some that we specifically use with great frequency are these (and I say loud **THANK YOU** to the authors!):

- R for Data Science
- An Introduction to Statistical Learning with Applications in R
- R: A self-learn tutorial
- Data Science in a Box
- stackoverflow.com, for example

Chapter 2

Syllabus

Instructor: Tobin Turner

Office Hours: mutually convenient time arranged by email e-mail: jttturner@presby.edu

2.1 Course Objectives and Learning Outcomes

This course is designed to introduce to data science. Students will apply statistical knowledge and techniques to both business and non-business contexts.

At the end of this course students should be able to:

- Demonstrate mastery of the statistical software in R and the Rstudio IDE.
- Data wrangle (the process of cleaning and unifying messy and complex data sets for easy access and analysis)
- Demonstrate mastery of single and multiple regression.
- Demonstrate mastery of these dplyr functions: filter, select, mutate, group_by, summarize, and tally.
- Demonstrate mastery of how business analytics is related to other business functions and is important to the success of the business entity.

This course will be focused on both understanding and applying key business analytical concepts. Although the text serves as a useful foundation for the concepts covered in the class, simple memorization of the material in the text will not be sufficient. Class participation, discussion, and application are critical.

2.2 Text and Resources

- The course website (primary resource)
- An Introduction to Statistical Learning with Applications in R; by Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani
- R: A self-learn tutorial.
- Other free, publicly available datasets and publications.

2.3 Performance Evaluation (Grading)

- Quizzes and Assignments - 20%
- Exam 1 - 20%
- Exam 2 - 20%
- Exam 3 - 20%
- Final Exam - 20%

2.3.1 Exams

Exams will cover assigned chapters in the textbook, other assigned readings, lectures, class exercises, class discussions, videos, and guest speakers. I will typically allocate time prior to each exam to clearly identify the body of knowledge each test will cover and to answer questions about the format and objectives of the exam.

2.3.2 Quizzes – DON'T MISS CLASS

- The average of all quizzes and assignments will comprise the Quizzes and Assignments - 20% portion of your final grade
- Quizzes are designed to prepare you for your exams and to ensure you stay up with the course material
- Missed quizzes cannot be made up later. Be present.

2.3.3 Final Average

- Final Average Grade
 - 90-100 A
 - 88-89 B+
 - 82-87 B
 - 80-81 B-
 - 78-79 C+
 - 72-77 C

- 70-71 C-
- 60-69 D
- 59 and below F

2.4 Class Participation:

I will frequently give readings or assignments for you to complete prior to the next class meeting. I expect you to fully engage the material: answer questions, pose questions, provide insightful observations. Keep in mind that quality is an important component in “participation.” Periodic cold calls will take place. I will also put students in the “hot seat” on occasion. In these class sessions, I may select a random group of students to lead us in the discussion and debate. Because the selection of participants will not be announced until class begins, everyone will be expected to prepare for the discussion. Reading the assigned chapters and articles are the best way to prepare for the discussion. If you have concerns about being called on in class, please see me to discuss. The purpose of the “hot seat” is not to stress or embarrass students, but to encourage students to actively engage the material.

2.5 Phones

Phones are not allowed to be used in class without the instructor’s prior consent. If you have a need of a phone during class please let me know before class. Unauthorized use of electronic devices may result in the lowering of the grade or dismissal from the class. **I mean this.**

The phone thing? I mean this.

2.6 Attendance

You are expected to be regular and punctual in your class attendance. Students are responsible for all the material missed and homework assignments made. If class is missed, notes/homework should be obtained from another student. If I am more than 15 minutes late, class is considered cancelled. No more than 4 absences are allowed during a semester. Exceeding the absence policy may result in receiving an F for the course. The professors roll is the official roll and students not present when roll is taken will be counted as absent. If a student must miss an exam, she or he must work out an agreeable time with the instructor to take the test prior to the exam being given. If a student misses a test due to an emergency, the student must inform the instructor as soon as is possible. In special cases, the instructor may allow the student to take a make-up exam.

Additionally, it is the student's responsibility to give the instructor one week's notice prior to each instance where accommodation will be required.

2.7 Honor Code and Plagiarism:

All assignments/exams must be your own work. Any copying or use of unauthorized assistance will be treated as a violation of PC's Honor Code. If you are unsure of what resources are allowed, please ask. Please note that all text longer than 7 words taken from ANY other source must be placed in quotations and cited. Also, summarizing ANY other source must also be cited. Using ANY other source and showing work to be your own is a violation of plagiarism and the honor code.

2.8 First-Generation Version:

I am a Presby First+ Advocate. I am here to support our current first-generation students. At Presbyterian College, first-generation students are those in which neither parent nor legal guardian graduated from a four-year higher education institution with a bachelor's degree. If you are a first-generation college student, please contact me. For more information about support for first-generation college students on our campus visit our Presby First+ webpage.

2.9 Continuing Advocate Version

I am a Presby First+ Advocate. I am committed to supporting first-generation students at Presbyterian College. At Presbyterian College, first-generation students are those in which neither parent nor legal guardian graduated from a four-year higher education institution with a bachelor's degree. If you are a first-generation college student, please contact me anytime or visit me during my office hours. For more information about support for first-generation college students on our campus visit our Presby First+ webpage.

2.10 Justice, Equity, Diversity, and Inclusion Syllabus Statements

Presbyterian College is committed to ensuring students from all faiths, identities and backgrounds have access to the services they need to thrive during their time learning in the classroom, living in our housing facilities, eating in the dining hall, or attending events on campus.

2.11 Accessible Education

To seek accommodations under the American Disabilities Act or Section 504 of the Rehabilitation Act, you must start the process yourself by completing the online form available on our website at <https://www.presby.edu/academics/academic-resources-support/accommodations-for-disabilities/> and submitting appropriate documentation to accommodations@presby.edu. Accommodations will become effective once the structured exchange and documentation review processes are complete.

2.12 Bias Reporting

Students are encouraged to report all acts of discrimination, spoken or written hate speech, harassment, offensive visual representations, threats, vandalism, and physical confrontations by using the bias reporting form available at <https://www.presby.edu/about/office-of-diversity-and-inclusion/bias-incident-reporting/>. All forms are sent directly to the VP for JEDI and will be addressed according to institutional policies and procedures.

2.13 Counseling and Health

Recognizing physical and mental health are vitally important to academic success, we encourage students to take advantage of the counseling and health care services available on campus. To schedule an appointment with Counseling, feel free to call (864) 833-8263 or email counselingandhealth@presby.edu.

We also have a nurse practitioner in the Student Health Center Monday through Friday 11am – 1:15pm. To schedule an appointment, call (864) 833-8400. If an appointment is needed outside of those hours and off-campus, please call Family Health Care at (864) 833-5986.

2.14 Presby First+ Program

The Presby First+ Program takes a holistic approach to helping first-generation college students accomplish their education and career goals. By connecting students with campus and community resources, the program positions students for success.

As a Presby First+ Advocate, I am committed to supporting first-generation college students. (Insert if applicable: I was a first-generation college student, meaning neither one of my parents had earned a bachelor's degree, so I may have experienced some of the things you will face during your journey.) Know

that you have lots of people wanting to see you thrive, so please use them as a resource. You can contact me anytime via email or stopping by my office. More information about the Presby First+ program is available on our website.

2.15 Title IX Syllabus Statement

In accordance with Title IX of the Education Amendments of 1972, Presbyterian College prohibits discrimination and harassment on the basis of sex. Presbyterian College strongly encourages the prompt reporting of, and is committed to a timely and fair resolution of, complaints of sex discrimination and harassment by students, faculty, staff, and third-party contractors. Sexual harassment includes quid pro quo harassment, sexual assault, dating violence, domestic violence, and stalking. Incidents of sex discrimination should be reported to the Title IX Coordinator, as outlined in the Sexual Misconduct Policy. The policy can be found on the Title IX page at [Title IX | Presbyterian College | Clinton, SC](#).

Chapter 3

Schedule

This is a tentative schedule, **BUT** I will do my very best to stick to it, so that you may plan accordingly!

Spring 2023

Date	Topic
Monday, January 9, 2023	First Day Fun!
Wednesday, January 11, 2023	R basics and install
Friday, January 13, 2023	R basics and workflows
Monday, January 16, 2023	Martin Luther King, Jr. Holiday
Wednesday, January 18, 2023	Objects and Arithmetic
Friday, January 20, 2023	QUIZ
Monday, January 23, 2023	Summaries and Subscribing
Wednesday, January 25, 2023	Case: GDH Ice Cream
Friday, January 27, 2023	QUIZ
Monday, January 30, 2023	Matrices and <code>mtcars</code>
Wednesday, February 1, 2023	Real Data
Friday, February 3, 2023	QUIZ
Monday, February 6, 2023	Review
Wednesday, February 8, 2023	EXAM 1
Friday, February 10, 2023	Exam 1 Review
Monday, February 13, 2023	Linear Regression
Wednesday, February 15, 2023	Multivariate Regression
Friday, February 17, 2023	Quiz
Monday, February 20, 2023	Social Dilemma and Review
Wednesday, February 22, 2023	<code>dplyr</code>
Friday, February 24, 2023	<code>dplyr</code>

Date	Topic
Monday, February 27, 2023	dplyr
Wednesday, March 1, 2023	Real data
Friday, March 3, 2023	Real data
Monday, March 6, 2023	Review
Wednesday, March 8, 2023	EXAM 2
Friday, March 10, 2023	Strategic Plan Launch (No classes)
Monday, March 13, 2023	SPRING BREAK
Wednesday, March 15, 2023	SPRING BREAK
Friday, March 17, 2023	SPRING BREAK
Monday, March 20, 2023	NFL
Wednesday, March 22, 2023	NFL
Friday, March 24, 2023	NFL
Monday, March 27, 2023	ggplot2
Wednesday, March 29, 2023	QUIZ
Friday, March 31, 2023	ggplot2
Monday, April 3, 2023	Advising Week
Wednesday, April 5, 2023	Advising Week
Friday, April 7, 2023	Easter Holidays
Monday, April 10, 2023	Easter Holidays
Wednesday, April 12, 2023	EXAM 3
Friday, April 14, 2023	Registration Week
Monday, April 17, 2023	
Wednesday, April 19, 2023	
Friday, April 21, 2023	
Monday, April 24, 2023	
Wednesday, April 26, 2023	
Friday, April 28, 2023	LAST DAY OF CLASSES
Wednesday, May 3, 2023	FINAL 1:30 p.m. – E period

Chapter 4

R basics and workflows

4.1 Basics of working with R at the command line and RStudio goodies

Launch RStudio/R.

You will first install R and then, RStudio.

- Installing R
- Installing RStudio
- Customizing RStudio
- RStudio Quick keys

4.2 In Rstudio - where we will live

Notice the default panes:

- Console (entire left)
- Environment/History (tabbed in upper right)
- Files/Plots/Packages/Help (tabbed in lower right)

4.3

Rstudio Console

FYI: you can change the default location of the panes, among many other things: Customizing RStudio.

Go into the Console, where we interact with the live R process.

You can make an object by assigning a value or statement to a letter or string. We use `<-` to assign objects meaning. Create and inspect the following object:

Your first analysis in R:

```
x <- 3 * 4
x
#> [1] 12
```

All R statements where you create objects – “assignments” – have this form:

```
objectName <- value
```

and in my head I hear, e.g., “x gets 12”.

You will make lots of assignments and the operator `<-` is a pain to type. Don’t be lazy and use `=`, although it would work, because it will just sow confusion later. Instead, utilize RStudio’s keyboard shortcut: `Alt + -` (the minus sign).

I want to be your friend. As a friend, I implore you, learn this:

In RStudio insert the `<-` assignment operator with `Option + -` (the minus sign) on a Mac, or `Alt + -` (the minus sign) on Windows.

Notice that RStudio automatically surrounds `<-` with spaces, which demonstrates a useful code formatting practice. Code is miserable to read on a good day. Give your eyes a break and use spaces.

RStudio offers many handy RStudio Quick keys. Also, `Alt+Shift+K` brings up a keyboard shortcut reference card.

Object names cannot start with a digit and cannot contain certain other characters such as a comma or a space. You will be wise to adopt a convention for demarcating words in names, but note that best practice is to choose ONE convention and stay true to it throughout your code.

```
i_use_snake_case
other.people.use.periods
evenOthersUseCamelCase
```

Make another assignment:

```
this_is_a_really_long_name <- 2.5
```

To inspect this, try out RStudio’s completion facility: type the first few characters, press TAB, add characters until you disambiguate, then press return.

Make another assignment:

```
turner_rocks <- 2 ^ 3
```

When making assignments, it is best practice to keep the names brief, yet descriptive. For instance, while the name “this_is_a_really_long_name” is accurate, so is “long_name” and this is much more intuitive and easy to read/type over and over.

Let’s try to inspect:

```
turnerrocks
#> Error in eval(expr, envir, enclos): object 'turnerrocks' not found
Turner_rocks
#> Error in eval(expr, envir, enclos): object 'Turner_rocks' not found
turner_rocks
#> [1] 8
```

Implicit contract with the computer / scripting language: Computer will do tedious computation for you. In return, you will be completely precise in your instructions. Typos matter. Case matters. Get better at typing.

R has a mind-blowing collection of built-in functions that are accessed like so:

```
functionName(arg1 = val1, arg2 = val2, and so on)
```

Let’s try using `seq()` which makes regular sequences of numbers and, while we’re at it, demo more helpful features of RStudio.

Type `se` and hit TAB. A pop up shows you possible completions. Specify `seq()` by typing more to disambiguate or using the up/down arrows to select. Notice the floating tool-tip-type help that pops up, reminding you of a function’s arguments. If you want even more help, press F1 as directed to get the full documentation in the help tab of the lower right pane. Now open the parentheses and notice the automatic addition of the closing parenthesis and the placement of cursor in the middle. Type the arguments 1, 10 and hit return. RStudio also exits the parenthetical expression for you. IDEs are great.

```
seq(1, 10)
#> [1] 1 2 3 4 5 6 7 8 9 10
```

The above also demonstrates something about how R resolves function arguments. You can always specify in `name = value` form. But if you do not, R attempts to resolve by position. So above, it is assumed that we want a sequence `from = 1` that goes `to = 10`. Since we didn't specify step size, the default value of `by` in the function definition is used, which ends up being 1 in this case. For functions I call often, I might use this resolve by position for the first argument or maybe the first two. After that, I always use `name = value`.

Make this assignment and notice similar help with quotation marks.

```
yo <- "hello world"
```

If you just make an assignment, you don't get to see the value, so then you're tempted to immediately inspect.

```
y <- seq(1, 10)
y
#> [1] 1 2 3 4 5 6 7 8 9 10
```

This common action can be shortened by surrounding the assignment with parentheses, which causes assignment and “print to screen” to happen.

It is best practice to always attempt to “print” your assignments after creating them. This will help alleviate the issue of searching 200+ lines of code for that one error causing argument.

```
(y <- seq(1, 10))
#> [1] 1 2 3 4 5 6 7 8 9 10
```

Not all functions have (or require) arguments:

```
date()
#> [1] "Wed Dec 14 10:08:21 2022"
```

Now look at your workspace – in the upper right pane. The workspace is where user-defined objects accumulate. You can also get a listing of these objects with commands:

```
objects()
#> [1] "this_is_a_really_long_name"
#> [2] "turner_rocks"
```

```
#> [3] "x"
#> [4] "y"
#> [5] "yo"
ls()
#> [1] "this_is_a_really_long_name"
#> [2] "turner_rocks"
#> [3] "x"
#> [4] "y"
#> [5] "yo"
```

If you want to remove the object named `y`, you can do this:

```
rm(y)
```

To remove everything:

```
rm(list = ls())
```

or click the broom in RStudio’s Environment pane.

4.4 Workspace and working directory

One day you will need to quit R, go do something else and return to your analysis later (this is a very happy day).

One day you will have multiple analyses going that use R and you want to keep them separate (a not so happy day).

One day you will need to bring data from the outside world into R and send numerical results and figures from R back out into the world (the happiest of days).

To handle these real life situations, you need to make two decisions:

- What about your analysis is “real”, i.e. will you save it as your lasting record of what happened?
- Where does your analysis “live”?

4.4.1 Workspace, `.RData`

As a beginning R user, it’s OK to consider your workspace “real”. *Very soon*, I urge you to evolve to the next level, where you consider your saved R scripts as “real”. (In either case, of course the input data is very much real and requires

preservation!) With the input data and the R code you used, you can reproduce *everything*. You can make your analysis fancier. You can get to the bottom of puzzling results and discover and fix bugs in your code. You can reuse the code to conduct similar analyses in new projects. You can remake a figure with different aspect ratio or save it as TIFF instead of PDF. You are ready to take questions. You are ready for the future.

If you regard your workspace as “real” (saving and reloading all the time), if you need to redo analysis ... you’re going to either redo a lot of typing (making mistakes all the way) or will have to mine your R history for the commands you used. Rather than [becoming an expert on managing the R history][rstudio-command-history], a better use of your time and psychic energy is to keep your “good” R code in a script for future reuse.

Because it can be useful sometimes, note the commands you’ve recently run appear in the History pane.

But you don’t have to choose right now and the two strategies are not incompatible. Let’s demo the save / reload the workspace approach.

Upon quitting R, you have to decide if you want to save your workspace, for potential restoration the next time you launch R. Depending on your set up, R or your IDE, e.g. RStudio, will probably prompt you to make this decision.

Quit R/RStudio, either from the menu, using a keyboard shortcut, or by typing `q()` in the Console. You’ll get a prompt like this:

```
Save workspace image to ~/.Rdata?
```

Note where the workspace image is to be saved and then click “Save”.

Using your favorite method, visit the directory where image was saved and verify there is a file named `.RData`. You will also see a file `.Rhistory`, holding the commands submitted in your recent session.

Restart RStudio. In the Console you will see a line like this:

```
[Workspace loaded from ~/.RData]
```

indicating that your workspace has been restored. Look in the Workspace pane and you’ll see the same objects as before. In the History tab of the same pane, you should also see your command history. You’re back in business. This way of starting and stopping analytical work will not serve you well for long but it’s a start.

4.4.2 Working directory

Any process running on your computer has a notion of its “working directory”. In R, this is where R will look, by default, for files you ask it to load. It also where, by default, any files you write to disk will go. Chances are your current working directory is the directory we inspected above, i.e. the one where RStudio wanted to save the workspace.

You can explicitly check your working directory with:

```
getwd()
```

It is also displayed at the top of the RStudio console.

As a beginning R user, it’s OK let your home directory or any other weird directory on your computer be R’s working directory. *Very soon*, I urge you to evolve to the next level, where you organize your analytical projects into directories and, when working on project A, set R’s working directory to the associated directory.

Although I do not recommend it, in case you’re curious, you can set R’s working directory at the command line like so:

```
setwd("~/myCoolProject")
```

Although I do not recommend it, you can also use RStudio’s Files pane to navigate to a directory and then set it as working directory from the menu: *Session > Set Working Directory > To Files Pane Location*. (You’ll see even more options there). Or within the Files pane, choose “More” and “Set As Working Directory”.

But there’s a better way. A way that also puts you on the path to managing your R work like an expert.

4.5 Exercises

1. Create an object called “cool_object” and assign it the number 100.
2. Create a new object called “big_brain” and multiply the object from question one by 15.
3. Print both objects.
4. Use base R functions to return today’s date and print it.
5. Create a sequence of numbers counting from 10 to 100 by 2.

6. Identify your working directory. What is it? Change it to where you want it.
7. Save the R script that answers questions 1 through 5 above. Save it; clean and close Rstudio; reopen your script and run it. Make sense?

Chapter 5

Objects and Arithmetic

5.1 Introduction

R stores information and operates on objects. The simplest objects are scalars, vectors and matrices. But there are many others: lists and dataframes for example. In advanced use of R it can also be useful to define new types of object, specific for particular application. We will stick with just the most commonly used objects here. An important feature of R is that it will do different things on different types of objects. For example, type:

```
4+6  
#> [1] 10
```

So, R does scalar arithmetic returning the scalar value 10. (In actual fact, R returns a vector of length 1 - hence the [1] denoting first element of the vector. We can assign objects values for subsequent use. For example:

```
x<-6  
y<-4  
z<-x+y
```

would do the same calculation as above, storing the result in an object called z. We can look at the contents of the object by simply typing its name:

```
z  
#> [1] 10
```

Storing things such as calculations as objects is extremely useful, especially in longer scripts. Mainly because you will likely call the same equation multiple times and if you need to revise it in any way you only need to change the initial assignment rather than every line.

5.2 Basic Functions

At any time we can list the objects which we have created:

```
ls()  
#> [1] "x" "y" "z"
```

Notice that `ls` is actually an object itself. Typing `ls` would result in a display of the contents of this object, in this case, the commands of the function. The use of parentheses, `ls()`, ensures that the function is executed and its result - in this case, a list of the objects in the directory - displayed. More commonly a function will operate on an object, for example:

```
sqrt(16)  
#> [1] 4
```

calculates the square root of 16. Objects can be removed from the current workspace with the `rm` function:

```
rm(x,y)
```

for example.

There are many standard functions available in R, and it is also possible to create new ones. Vectors can be created in R in a number of ways. We can describe all of the elements:

```
z<-c(5,9,1,0)
```

Note the use of the function `c` to concatenate or glue together individual elements. This function can be used much more widely, for example

```
x<-c(5,9)  
y<-c(1,0)  
z<-c(x,y)
```

would lead to the same result by gluing together two vectors to create a single vector.

Sequences can be generated as follows:

```
x<-1:10
```

while more general sequences can be generated using the `seq` command. For example:


```
x<-c(6,8,9)
y<-c(1,2,4)
x+y
#> [1]  7 10 13
```

and

```
x*y
#> [1]  6 16 36
```

showing that R uses componentwise arithmetic on vectors. R will also try to make sense if objects are mixed. For example:

```
x<-c(6,8,9)
x+2
#> [1]  8 10 11
```

though care should be taken to make sure that R is doing what you would like it to in these circumstances.

Two particularly useful functions worth remembering are `length` which returns the length of a vector (i.e. the number of elements it contains) and `sum` which calculates the sum of the elements of a vector.

5.3 Statistics and Summaries

One of the most useful functions of R is its ability to perform statistical analysis on large pieces of data. Later in this book we will cover how this analysis can be used to build complex models and test their accuracy. For now, the focus will be on how to perform basic statistical analysis and the definitions of the more basic terms.

5.3.1 Statistical Analysis

You've undoubtedly have heard the terms mean, median, standard deviation and variance. However, if asked to define these terms and provide examples of their usefulness, what would you say? This section is going to answer the first part of that question.

Mean is a pretty straight forward concept for most. Commonly referred to as the “average”, you find this number by adding all the numbers in your data set together and then divide by the total number of points you used. This is easily done on a calculator if you are working with less than 20 digits, however this

is almost never the case in data science. Often, especially with R, you will be dealing with data sets with hundreds, thousands, or maybe even hundreds of thousands of data points. Fortunately, R takes the pain away from this process and offers a very simple function that will do this for you. Start by creating an object that contains a list of numbers, and then simply use the `mean` function to calculate the average of your variable.

```
a <- c(3.8,4,3.1,2.1,12.6,17,8.43,11,2,3,9,5,3,0.5)

mean(a)
#> [1] 6.037857
```

The median of a data set is data point that fall in the middle of all the other points when they are ordered from lowest to highest. Again, something easily found even without a calculator if you have small data sets, however humans can miscount and some data sets are just too massive for us to do on our own. R never makes these counting mistakes and by using a simple function you save yourself hours of mundane counting. Start the same way you found mean by creating an object with various numbers and then simply use the `median` function.

```
b <- c(3.8,4,3.1,2.1,12.6,17,8.43,11,2,3,9,5,3,0.5)

median(b)
#> [1] 3.9
```

Many people have heard of standard deviation, but not many can provide an accurate definition. Standard deviation is the measure of variance (see next paragraph) between numbers in a data set and that data set's mean. Typically a lower standard deviation is what you want to see as this proves there is little variance in your data set and that typically means a higher correlation. If you wanted to find this by hand you would have to find the square root of the variance between each point and the mean. This is easy if you have all the numbers you need, but that is almost never the case and thus you would have to calculate all that variance by hand and that is too much work. Why take all those extra steps when R can do it for you? The `sd()` function does all the work without any hassle, all you have to do is identify the object you want it to find the standard deviation of.

```
d <- c(3.8,4,3.1,2.1,12.6,17,8.43,11,2,3,9,5,3,0.5)

sd(d)
#> [1] 4.821066
```

Finally, we can cover variance. Variance is the measure of distance between two numbers. This statistic is often used to measure how spread out your data is and

can be used to determine your model's accuracy. Typically, a higher variance means your model is inaccurate. So, when building models and reviewing your summary statistics, you want your variance to be as low as possible. There is a `var()` function that will find the variance of a variable for you, however typically you will create a confusion matrix (discussed in a later section) that will help you determine the variance, and thus accuracy of your model.

5.4 Exercises

1. Define

```
x<-c(4,2,6)
y<-c(1,0,-1)
```

2. Decide what the result will be of the following:

- `length(x)`
- `sum(x)`
- `sum(x^2)`
- `x+y`
- `x*y`
- `x-2`
- `x^2`

3. Use R to check your answers.

4. Decide what the following sequences are and use R to check your answers:

- `7:11`
- `seq(2,9)`
- `seq(4,10,by=2)`
- `seq(3,30,length=10)`
- `seq(6,-4,by=-2)`

5. Determine what the result will be of the following R expressions, and then use R to check if you are right:

- `rep(2,4)`
- `rep(c(1,2),4)`
- `rep(c(1,2),c(4,4))`
- `rep(1:4,4)`
- `rep(1:4,rep(3,4))`

6. Use the `rep` function to define simply the following vectors in R.

- `6,6,6,6,6,6`
- `5,8,5,8,5,8,5,8`
- `5,5,5,5,8,8,8,8`

Chapter 6

Summaries and Subscripting

6.1 Introduction

Let's suppose we've collected some data from an experiment and stored them in an object `x`:

```
x<-c(7.5,8.2,3.1,5.6,8.2,9.3,6.5,7.0,9.3,1.2,14.5,6.2)
```

Some simple summary statistics of these data can be produced:

```
mean(x)
#> [1] 7.216667
var(x)
#> [1] 11.00879
sd(x)
#> [1] 3.317949
summary(x)
#>   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
#>  1.200  6.050   7.250   7.217   8.475  14.500
```

which should all be self explanatory.

It may be, however, that we subsequently learn that the first 6 data points correspond to measurements made on one machine, and the second six on another machine.

This might lead us to want to summarize the two sets of data separately, so we would need to extract from `x` the two relevant subvectors. This is achieved by subscripting:

```
x[1:6]
#> [1] 7.5 8.2 3.1 5.6 8.2 9.3
x[7:12]
#> [1] 6.5 7.0 9.3 1.2 14.5 6.2
summary(x[1:6])
#>      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
#>  3.100   6.075   7.850   6.983   8.200   9.300
summary(x[7:12])
#>      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
#>  1.200   6.275   6.750   7.450   8.725  14.500
```

Other subsets can be created in the obvious way. For example:

```
x[c(2,4,9)]
#> [1] 8.2 5.6 9.3
```

Negative integers can be used to exclude particular elements. For example

```
x[-(1:6)]
#> [1] 6.5 7.0 9.3 1.2 14.5 6.2
```

has the same effect as

```
x[7:12]
#> [1] 6.5 7.0 9.3 1.2 14.5 6.2
```

6.2 Exercises (Summaries and Subscripting)

1. If `x<- c(5,9,2,3,4,6,7,0,8,12,2,9)` decide what each of the following is and use R to check your answers:
 - (a) `x[2]`
 - (b) `x[2:4]`
 - (c) `x[c(2,3,6)]`
 - (d) `x[c(1:5,10:12)]`
 - (e) `x[-(10:12)]`
2. The data `y<-c(33,44,29,16,25,45,33,19,54,22,21,49,11,24,56)` contain sales of milk in gallons for 5 days in three different shops (the first 3 values are for shops 1,2 and 3 on Monday, etc.) Produce a statistical summary of the sales for each day of the week and also for each shop.

Chapter 7

Matrices

7.1 CBind and RBind

Matrices can be created in R in a variety of ways. Perhaps the simplest is to create the columns (just a couple of objects) and then glue them together with the command `cbind`. For example,

```
x<-c(5,7,9)
y<-c(6,3,4)
z<-cbind(x,y)
z
#>      x y
#> [1,] 5 6
#> [2,] 7 3
#> [3,] 9 4
```

The dimension of a matrix can be checked with the `dim` command:

```
dim(z)
#> [1] 3 2
```

[1] 3 2 i.e., three rows and two columns. There is a similar command, `rbind`, for building matrices by gluing rows together.

The functions `cbind` and `rbind` can also be applied to matrices themselves (provided the dimensions match) to form larger matrices. For example,

```
rbind(z,z)
#>      x y
#> [1,] 5 6
#> [2,] 7 3
#> [3,] 9 4
#> [4,] 5 6
#> [5,] 7 3
#> [6,] 9 4
```

```
#> [1,] 5 6  
#> [2,] 7 3  
#> [3,] 9 4  
#> [4,] 5 6  
#> [5,] 7 3  
#> [6,] 9 4
```

7.1.1 Review Questions

- 1) Create a matrix made up of two columns showing the GPAs and number of hours studied by seven students.
- 2) Recreate the following matrix in R:

```
#>      x  y  
#> [1,] 5 3.4  
#> [2,] 7 4.0  
#> [3,] 2 2.5  
#> [4,] 3 3.2  
#> [5,] 8 2.8  
#> [6,] 4 3.1  
#> [7,] 2 3.6
```

- 3) Using the appropriate function, combine the two matrices you created above.

7.2 Matrix Function

Matrices can also be built by explicit construction via the function `matrix`. For example,

```
z<-matrix(c(5,7,9,6,3,4),nrow=3)
```

results in a matrix `z` identical to `z` above. Notice that the dimension of the matrix is determined by the size of the vector and the requirement that the number of rows is 3, as specified by the argument `nrow=3`. As an alternative we could have specified the number of columns with the argument `ncol=2` (obviously, it is unnecessary to give both). Notice that the matrix is ‘filled up’ column-wise. If instead you wish to fill up row-wise, add the option `byrow=T`. For example,

```
z<-matrix(c(5,7,9,6,3,4),nr=3,byrow=T)
z
#>      [,1] [,2]
#> [1,]    5    7
#> [2,]    9    6
#> [3,]    3    4
```

Notice that the argument `nrow` has been abbreviated to `nr`. Such abbreviations are always possible for function arguments provided it induces no ambiguity - if in doubt always use the full argument name.

As usual, R will try to interpret operations on matrices in a natural way. For example, with `z` as above, and

```
y<-matrix(c(1,3,0,9,5,-1),nrow=3,byrow=T)
y
#>      [,1] [,2]
#> [1,]    1    3
#> [2,]    0    9
#> [3,]    5   -1
```

we obtain

```
y+z
#>      [,1] [,2]
#> [1,]    6   10
#> [2,]    9   15
#> [3,]    8    3
```

and

```
y*z
#>      [,1] [,2]
#> [1,]    5   21
#> [2,]    0   54
#> [3,]   15  -4
```

Other useful functions on matrices are to transpose a matrix:

```
z
#>      [,1] [,2]
#> [1,]    5    7
#> [2,]    9    6
#> [3,]    3    4
```

```
t(z)
#>      [,1] [,2] [,3]
#> [1,]    5    9    3
#> [2,]    7    6    4
```

As with vectors it is useful to be able to extract sub-components of matrices. In this case, we may wish to pick out individual elements, rows or columns. As before, the `[]` notation is used to subscript. The following examples should make things clear:

```
z[1,1]
#> [1] 5
```

```
z[c(2,3),2]
#> [1] 6 4
```

```
z[,2]
#> [1] 7 6 4
```

```
z[1:2,]
#>      [,1] [,2]
#> [1,]    5    7
#> [2,]    9    6
```

So, in particular, it is necessary to specify which rows and columns are required, whilst omitting the integer for either dimension implies that every element in that dimension is selected.

7.3 Exercises

1. Create this matrix in R

```
#>      [,1] [,2] [,3] [,4] [,5]
#> [1,]    1    7    8    11   -5
#> [2,]    3    8    6     3   -9
#> [3,]    0   11   14   14   14
```

2. Create in R these matrices:

```

x
#>      [,1] [,2]
#> [1,]    1    7
#> [2,]    8   11
#> [3,]    5    9
y
#>      [,1] [,2]
#> [1,]    6    8
#> [2,]    2    1
#> [3,]    1   -7

```

3. Calculate the following and check your answers in R:

- (a) $2*x$
- (b) $x*x$
- (c) $t(y)$

```

#>      [,1] [,2]
#> [1,]    2   14
#> [2,]   16   22
#> [3,]   10   18
#>      [,1] [,2]
#> [1,]    1   49
#> [2,]   64  121
#> [3,]   25   81
#>      [,1] [,2] [,3]
#> [1,]    6    2    1
#> [2,]    8    1   -7

```

4. With x and y as above, calculate the effect of the following subscript operations and check your answers in R.

- (a) $x[1,]$
- (b) $x[2,]$
- (c) $x[,2]$
- (d) $y[1,2]$
- (e) $y[,2:3]$

Chapter 8

Data Frames

Data Frames are data displayed in a format as a table.

Data Frames can have different types of data inside it. While the first column can be character, the second and third can be numeric or logical. However, each column should have the same type of data.

Use the `data.frame()` function to create a data frame:

```
# Create a data frame
Data_Frame <- data.frame (
  Training = c("Strength", "Stamina", "Other"),
  Pulse = c(100, 150, 120),
  Duration = c(60, 30, 45)
)

# Print the data frame
Data_Frame
#>   Training Pulse Duration
#> 1 Strength   100      60
#> 2 Stamina   150      30
#> 3  Other    120      45
```

Adding and Editing a Data Frame

```
# Add a new row
New_row_DF <- rbind(Data_Frame, c("Strength", 110, 110))

# Print the new row
New_row_DF
#>   Training Pulse Duration
```

```

#> 1 Strength 100      60
#> 2 Stamina 150      30
#> 3 Other 120      45
#> 4 Strength 110     110
# Add a new column
New_col_DF <- cbind(Data_Frame, Steps = c(1000, 6000, 2000))

# Print the new column
New_col_DF
#> Training Pulse Duration Steps
#> 1 Strength 100      60 1000
#> 2 Stamina 150      30 6000
#> 3 Other 120      45 2000

# Combining Data Frames
# Use the rbind() function to combine two or more data frames in R vertically:

Data_Frame1 <- data.frame (
  Training = c("Strength", "Stamina", "Other"),
  Pulse = c(100, 150, 120),
  Duration = c(60, 30, 45)
)

Data_Frame2 <- data.frame (
  Training = c("Stamina", "Stamina", "Strength"),
  Pulse = c(140, 150, 160),
  Duration = c(30, 30, 20)
)

New_Data_Frame <- rbind(Data_Frame1, Data_Frame2)
New_Data_Frame
#> Training Pulse Duration
#> 1 Strength 100      60
#> 2 Stamina 150      30
#> 3 Other 120      45
#> 4 Stamina 140      30
#> 5 Stamina 150      30
#> 6 Strength 160      20

```

8.1 Exercises

1. Write a R program to create the following data frame:


```
# Create the data frame.
emp.data <- data.frame(
  emp_id = c (1:5),
  emp_name = c("Rick","Dan","Michelle","Ryan","Gary"),
  salary = c(623.3,515.2,611.0,729.0,843.25),
  start_date = as.Date(c("2012-01-01", "2013-09-23", "2014-11-15", "2014-05-11",
    "2015-03-27")),
  stringsAsFactors = FALSE
)
# Print the data frame.
print(emp.data)
#>   emp_id emp_name salary start_date
#> 1      1    Rick 623.30 2012-01-01
#> 2      2     Dan 515.20 2013-09-23
#> 3      3 Michelle 611.00 2014-11-15
#> 4      4     Ryan 729.00 2014-05-11
#> 5      5     Gary 843.25 2015-03-27
```


Chapter 9

Preloaded data and mtcars

R comes with several built-in data sets, **which are generally used as demo data for playing with R functions.**

To see the datasets type:

```
data()
```

9.1 Practicing with mtcars data set

This demonstration is based on the dataset `mtcars`.

1. Read in `mtcars`

```
data(mtcars)
```

2. View first few rows and last few rows of `mtcars` dataframe using functions `head()` and `tail()`

```
head(mtcars)
```

```
#>           mpg cyl  disp  hp  drat   wt  qsec vs am
#> Mazda RX4      21.0   6  160 110 3.90 2.620 16.46 0  1
#> Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02 0  1
#> Datsun 710      22.8   4  108  93 3.85 2.320 18.61 1  1
#> Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44 1  0
#> Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02 0  0
#> Valiant         18.1   6  225 105 2.76 3.460 20.22 1  0
#>           gear carb
#> Mazda RX4         4   4
```

```
#> Mazda RX4 Wag      4      4
#> Datsun 710          4      1
#> Hornet 4 Drive     3      1
#> Hornet Sportabout  3      2
#> Valiant             3      1
tail(mtcars)
#>      mpg  cyl  disp  hp drat   wt  qsec vs am
#> Porsche 914-2  26.0   4 120.3  91 4.43 2.140 16.7 0 1
#> Lotus Europa   30.4   4  95.1 113 3.77 1.513 16.9 1 1
#> Ford Pantera L  15.8   8 351.0 264 4.22 3.170 14.5 0 1
#> Ferrari Dino   19.7   6 145.0 175 3.62 2.770 15.5 0 1
#> Maserati Bora   15.0   8 301.0 335 3.54 3.570 14.6 0 1
#> Volvo 142E      21.4   4 121.0 109 4.11 2.780 18.6 1 1
#>      gear carb
#> Porsche 914-2    5    2
#> Lotus Europa     5    2
#> Ford Pantera L   5    4
#> Ferrari Dino     5    6
#> Maserati Bora    5    8
#> Volvo 142E       4    2
```

3. Some info about mtcars dataframe using function `colnames()`, `rownames()`, `summary()` and `dim()`

```
colnames(mtcars)
#> [1] "mpg" "cyl" "disp" "hp" "drat" "wt" "qsec" "vs"
#> [9] "am" "gear" "carb"
rownames(mtcars)
#> [1] "Mazda RX4" "Mazda RX4 Wag"
#> [3] "Datsun 710" "Hornet 4 Drive"
#> [5] "Hornet Sportabout" "Valiant"
#> [7] "Duster 360" "Merc 240D"
#> [9] "Merc 230" "Merc 280"
#> [11] "Merc 280C" "Merc 450SE"
#> [13] "Merc 450SL" "Merc 450SLC"
#> [15] "Cadillac Fleetwood" "Lincoln Continental"
#> [17] "Chrysler Imperial" "Fiat 128"
#> [19] "Honda Civic" "Toyota Corolla"
#> [21] "Toyota Corona" "Dodge Challenger"
#> [23] "AMC Javelin" "Camaro Z28"
#> [25] "Pontiac Firebird" "Fiat X1-9"
#> [27] "Porsche 914-2" "Lotus Europa"
#> [29] "Ford Pantera L" "Ferrari Dino"
#> [31] "Maserati Bora" "Volvo 142E"
summary(mtcars)
```

```

#>      mpg      cyl      disp
#> Min.   :10.40   Min.   :4.000   Min.    : 71.1
#> 1st Qu.:15.43   1st Qu.:4.000   1st Qu.:120.8
#> Median :19.20   Median :6.000   Median :196.3
#> Mean   :20.09   Mean   :6.188   Mean   :230.7
#> 3rd Qu.:22.80   3rd Qu.:8.000   3rd Qu.:326.0
#> Max.   :33.90   Max.   :8.000   Max.   :472.0
#>      hp      drat      wt
#> Min.   : 52.0   Min.   :2.760   Min.   :1.513
#> 1st Qu.: 96.5   1st Qu.:3.080   1st Qu.:2.581
#> Median :123.0   Median :3.695   Median :3.325
#> Mean   :146.7   Mean   :3.597   Mean   :3.217
#> 3rd Qu.:180.0   3rd Qu.:3.920   3rd Qu.:3.610
#> Max.   :335.0   Max.   :4.930   Max.   :5.424
#>      qsec      vs      am
#> Min.   :14.50   Min.   :0.0000   Min.   :0.0000
#> 1st Qu.:16.89   1st Qu.:0.0000   1st Qu.:0.0000
#> Median :17.71   Median :0.0000   Median :0.0000
#> Mean   :17.85   Mean   :0.4375   Mean   :0.4062
#> 3rd Qu.:18.90   3rd Qu.:1.0000   3rd Qu.:1.0000
#> Max.   :22.90   Max.   :1.0000   Max.   :1.0000
#>      gear      carb
#> Min.   :3.000   Min.   :1.000
#> 1st Qu.:3.000   1st Qu.:2.000
#> Median :4.000   Median :2.000
#> Mean   :3.688   Mean   :2.812
#> 3rd Qu.:4.000   3rd Qu.:4.000
#> Max.   :5.000   Max.   :8.000
dim(mtcars)
#> [1] 32 11

```

4. To calculate the variance of weight:

```

var(mtcars$wt)
#> [1] 0.957379


```

5. To get the histogram of hp, the code below will produce a histogram:

```

hist(mtcars$hp)

```




6. To calculate the quantiles by percent:

```
quantile(mtcars$wt, c(.2, .4, .8))  
#> 20% 40% 80%  
#> 2.349 3.158 3.770
```

9.2 Excerises for you:

1. Find the minimum and maximum value of `mpg`
2. Find the mean and standard deviation of data variable `mpg`
3. What variable has a 3rd quartile value of 180.0?
4. Create and explain what this means



5. Create and explain what this means

```
#>           mpg           cyl           disp           hp
#> mpg  1.0000000 -0.8521620 -0.8475514 -0.7761684
#> cyl -0.8521620  1.0000000  0.9020329  0.8324475
#> disp -0.8475514  0.9020329  1.0000000  0.7909486
#> hp  -0.7761684  0.8324475  0.7909486  1.0000000
#> drat  0.6811719 -0.6999381 -0.7102139 -0.4487591
#> wt   -0.8676594  0.7824958  0.8879799  0.6587479
#> qsec  0.4186840 -0.5912421 -0.4336979 -0.7082234
#> vs   0.6640389 -0.8108118 -0.7104159 -0.7230967
#> am    0.5998324 -0.5226070 -0.5912270 -0.2432043
#> gear  0.4802848 -0.4926866 -0.5555692 -0.1257043
#> carb -0.5509251  0.5269883  0.3949769  0.7498125
#>           drat           wt           qsec           vs
#> mpg  0.68117191 -0.8676594  0.41868403  0.6640389
#> cyl -0.69993811  0.7824958 -0.59124207 -0.8108118
#> disp -0.71021393  0.8879799 -0.43369788 -0.7104159
#> hp  -0.44875912  0.6587479 -0.70822339 -0.7230967
#> drat 1.00000000 -0.7124406  0.09120476  0.4402785
#> wt   -0.71244065  1.0000000 -0.17471588 -0.5549157
#> qsec  0.09120476 -0.1747159  1.00000000  0.7445354
#> vs   0.44027846 -0.5549157  0.74453544  1.0000000
#> am    0.71271113 -0.6924953 -0.22986086  0.1683451
```

```
#> gear 0.69961013 -0.5832870 -0.21268223 0.2060233
#> carb -0.09078980 0.4276059 -0.65624923 -0.5696071
#>
#>          am          gear          carb
#> mpg 0.59983243 0.4802848 -0.55092507
#> cyl -0.52260705 -0.4926866 0.52698829
#> disp -0.59122704 -0.5555692 0.39497686
#> hp -0.24320426 -0.1257043 0.74981247
#> drat 0.71271113 0.6996101 -0.09078980
#> wt -0.69249526 -0.5832870 0.42760594
#> qsec -0.22986086 -0.2126822 -0.65624923
#> vs 0.16834512 0.2060233 -0.56960714
#> am 1.00000000 0.7940588 0.05753435
#> gear 0.79405876 1.0000000 0.27407284
#> carb 0.05753435 0.2740728 1.00000000
```

6. Create a variable called `efficiency` which is `mpg` divided by `weight`. Which car has the max `efficiency` and what is this value?
7. Which variable in this dataset has the greatest standard deviation?
8. How many cars have 3 gears?
9. How many cars get more than 17 mpg?

Chapter 10

More simple data wrangling

10.1 a nice, fun little matrix for you

```
#>      [,1] [,2] [,3] [,4] [,5]
#> [1,]    1    4    7   10   13
#> [2,]    2    5    8   11   14
#> [3,]    3    6    9   12   15
```

1. Write the code that creates this matrix:
2. Write DIFFERENT code that creates this matrix in an alternate way:
3. In the matrix above, what does [,4] mean?
4. What code would return the value in the 3rd column and 3rd row?
5. What single line of code would give you the average of all the numbers in columns 2, 4, and 5 and in rows 1 and 3?
6. turn `x` into a data frame.
7. How do you **know** you have turned `x` into a data frame?

10.2 More fun (this class is really awesome isn't it?)

```
df
#>      X1 X2 X3 X4 X5 X6 X7 X8 X9 X10
#> 1    1 11 21 31 41 51 61 71 81 91
#> 2    2 12 22 32 42 52 62 72 82 92
#> 3    3 13 23 33 43 53 63 73 83 93
#> 4    4 14 24 34 44 54 64 74 84 94
#> 5    5 15 25 35 45 55 65 75 85 95
#> 6    6 16 26 36 46 56 66 76 86 96
#> 7    7 17 27 37 47 57 67 77 87 97
#> 8    8 18 28 38 48 58 68 78 88 98
#> 9    9 19 29 39 49 59 69 79 89 99
#> 10 10 20 30 40 50 60 70 80 90 100
```

1. Consider the dataframe above called `df`. What would running this code return `sum(df[7,7:10])`
2. How can you tell if an object in R is a dataframe?
3. How could you create the dataframe above called `df`?
4. What code would return the average of row 2 of `df`?
5. Consider `mtcars` dataset that comes preloaded with R that looks like this:

```
head(mtcars)
#>      mpg  cyl  disp  hp  drat    wt   qsec vs  am
#> Mazda RX4      21.0   6  160 110 3.90 2.620 16.46 0  1
#> Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02 0  1
#> Datsun 710     22.8   4  108  93 3.85 2.320 18.61 1  1
#> Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44 1  0
#> Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02 0  0
#> Valiant        18.1   6  225 105 2.76 3.460 20.22 1  0
#>      gear carb
#> Mazda RX4      4    4
#> Mazda RX4 Wag  4    4
#> Datsun 710     4    1
#> Hornet 4 Drive  3    1
#> Hornet Sportabout 3    2
#> Valiant        3    1
```

6. Why do I get this error when I run the code below: `Error in plot(hp, mpg) : object 'hp' not found?`

```
plot(hp,mpg)
```

```
Error in plot(hp, mpg) : object 'hp' not found
```

10.2. *MORE FUN (THIS CLASS IS REALLY **AWESOME** ISN'T IT?)* 51

Bonus: What is a topic that you find confusing at this point in class?

Chapter 11

GDH Ice Cream

11.1 Problem Introduction

GDH provides ice cream for its wonderful customers. I LOVE GDH. Do you love it as much as me (let's discuss)?

In the last three years GDH used ice cream, in pounds, by month, as shown in the attached file.

```
#>  Month.Name year1 year2 year3
#>      Jan     60     67     64
#>      Feb     68     67     72
#>      Mar     83     62     61
#>      Apr    102     95    107
#>      May     95    105    101
#>      Jun     57     89     75
#>      Jul     61     57     81
#>      Aug    109    109    104
#>      Sep     56     86     88
#>      Oct     53     53     65
#>      Nov     74     72     72
#>      Dec     73     64     65
```

11.2 Assignment

Please answer the following questions *using R*.

GDH provides ice cream cones for its customers. In the last three years GDH used ice cream, in pounds, by month, as shown in the attached file.

1. In R, create the above data frame and name it `ice.cream`
2. What is another way you could have created the same data set?
3. Using R, what is the mean using for the months of Feb and Oct?
4. Create a chart showing ice cream use over time.
5. Which year used the most ice cream?
6. Which month has the highest standard deviation of ice cream use?
7. Which year has the highest standard deviation of ice cream use?
8. Also, you May want to check out this link to look at something called dataframes that may help with this assignment (but is not absolutely necessary) <https://www.rdocumentation.org/packages/base/versions/3.6.2/topics/data.frame>
9. Can you transpose your matrix?
10. Can you add meaningful row names and column names?


Chapter 12

Exam 1 Practice

You should expect to be able to do 7 to 10 of these type problems for your exam (50 minutes).

1. Load the `rock` data set. How many observations and variables are in the `rock` data set?
2. What was the purpose of this dataset?
3. What are the units of measurement for `peri`?
4. In the `rock` data set what is the standard deviation of `area`?
2683.8488617
5. How many rocks have an `area` of greater than 4333? **41**
6. How many rocks have an `area` of greater than 4333 and less than 7333?
14
7. Create a variable called `good.fun` which is area divided by perimeter squared. What is the average of `good.fun`? **0.0020362**
8. Add a new observation that to this dataframe that has an `area` of 17777, `peri` of 7777, `shape` of 7777, and a `perm` of 7777. Can you do this?
9. What is the mean of area for including the new rock you have added?
10. Using the `women` data set, what is the average BMI for all women?
11. What is the average BMI for `women`?
12. What is the correlation between height and weight in the `women` dataset?
13. Make a histogram plot of the area variable from `rock`.

14. Use a bivariate plot of `rock` data to identify which variable are most correlated. Confirm your answer with a correlation table. Were you right?
15. Do this math in R: 3333 squared divided by the square root of 78.
16. Make this sequence using `rep` and `seq` functions: **1 2 3 1 2 3 1 2 3 1 2 3**
17. Make this sequence using `rep` and `seq` functions: **97, 96, 95, 94, 93, 92, 91, 90, 89, 97, 96, 95, 94, 93, 92, 91, 90, 89, 97, 96, 95, 94, 93, 92, 91, 90, 89**



18. Make this plot:
19. What is the heaviest car in the `mtcars` dataset in kg?
20. How many cars in `mtcars` dataset have more than 4 cylinders?
21. Create this dataframe called `Gerry`:
22. Add column called `aa` to `Gerry` which are the numbers **11120, 11121, 11122, 11123, 11124, 11125, 11126, 11127, 11128, 11129**
23. Have fun. You are doing **awesome!**

#Machine Learning?

Have we talked about the Target Pregnancy Story yet?

How Target Figured Out A Teen Girl Was Pregnant Before Her Father Did

This is where the fun stuff begins! What we have learned up to this point has barely scratched the surface of what R is capable of. In the world of data science, R is used for three primary purposes, those purposes are **(1) data transformation, (2) data wrangling, (3) machine learning**. The other two purposes have been covered in earlier chapters of this book. The reason we covered the other topics first is that that lay the foundation. In the real world, it is likely you will never be given a clean data set, and you will have to do some wrangling and transformation before anything else is possible. After all, in the experience of many data science students, cleaning the data is the most tedious and time consuming process of a project.

Enough of the old stuff, what is machine learning? According to the Merriam-Webster Dictionary, machine learning is “the process by which a computer is able to improve its own performance by continuously incorporating new data into an existing statistical model.” Let’s take a trip back to the Target story discussed in the introductory chapter. The data scientist, or more likely data scientists (collaborative work is essential in this field), that worked on that model were likely experts in machine learning. They were able to train the computer to look through thousands (probably more) of customers’ data and the computer, based off the algorithms written by these “data nerds,” was able to predict whether a customer was pregnant! Think of other possible applications of this technology? We could predict how well a student will perform on an exam, the risk of someone suffering from a heart attack, or the likelihood that someone will default on a loan. Every field in existence today could find a way to implement machine learning to optimize their business.

There are two branches of machine learning, supervised and unsupervised. Both have their own unique uses, however in this course we will focus on supervised machine learning. Supervised machine learning required us to provide a clean data set with clearly defined variables and instructions. Essentially, we give the computer the information it needs and provide it with specific instructions detailing what we would like to see happen, and it does the rest. Linear regression is typically the first method of supervised learning people are introduced to, and it will be the focus of this chapter.


Note, these concepts are not all common sense and can be difficult to wrap your head around at times. Be sure to constantly turn to your instructor or peers for assistance and remember that there are hundreds of online resources at your disposal. As with anything though, practice makes perfect. The popular rule states that mastering a skill can take upwards of 10,000 hours! Now, this course is not going to take you 10 years to complete, however the goal is that by the end of this chapter you will know your way around the basics of linear regression.

Chapter 13


Quick Linear Regression

13.1 Quick Linear regression using Loblolly

1. load `Loblolly` and create a scatter plot of the data so plot so that age is the independent variable and height is the dependent variable.



2. Notice that R automatically labeled the x- and y-axes, but we also want our scatter plot to have a main title. To add a title, use the command `title(main = "Loblolly Pine Tree Heights")`.




3. To find a linear model that relates the age and height of the loblolly pine trees, we will use the command `fit1<-lm(Loblolly$height~Loblolly$age)`.

Think of `lm(Loblolly$height~Loblolly$age)` as the slope-intercept form ($y=mx+b$).

4. To see the model, type `fit1`

```
fit1 <- lm(Loblolly$height~Loblolly$age)
fit1
#>
#> Call:
#> lm(formula = Loblolly$height ~ Loblolly$age)
#>
#> Coefficients:
#> (Intercept)  Loblolly$age
#>      -1.312      2.591
```

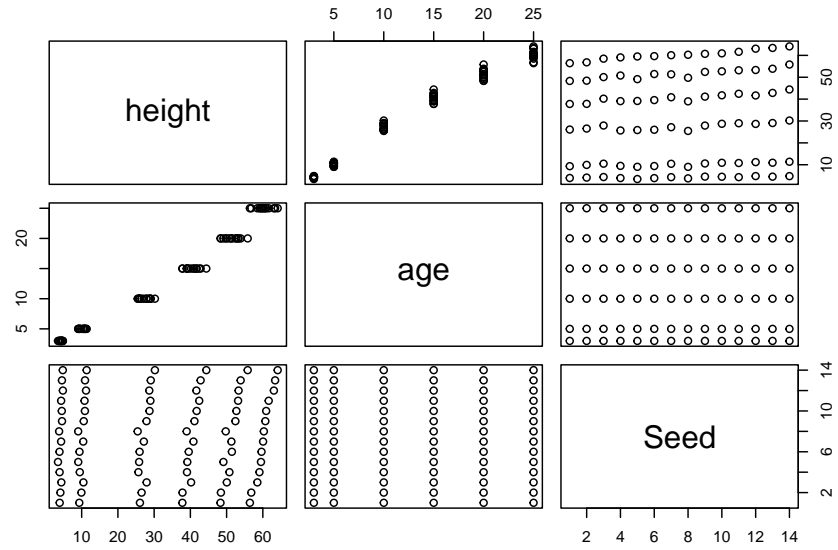
5. Now we want to add the graph of this line of best fit to our scatter plot. To do this, use the command `abline(fit1)`.



9. The final piece of information we want about our data is the correlation of the age and height of the Loblolly pine trees. To find the correlation coefficient, use the command `cor(Loblolly$height, Loblolly$age)`

```
#> [1] 0.9899132
#>
#> Pearson's product-moment correlation
#>
#> data:  Loblolly$height and Loblolly$age
#> t = 63.272, df = 82, p-value < 2.2e-16
#> alternative hypothesis: true correlation is not equal to 0
#> 95 percent confidence interval:
#>  0.9844505 0.9934631
#> sample estimates:
#>      cor
#> 0.9899132
```

10. What does this command do and mean: `plot(Loblolly)`?



Chapter 14

Linear Regression with `mtcars`

Remember: \sim here means “explained by”, so the formula `mpg ~ wt` means we are predicting `mpg` as explained by `wt`. The most helpful way to view the output is with:

14.1 Exercises for you

14.1.1 `mtcars`

1. Which variable in the `mtcars` dataset do you think best predicts `mpg` and why?
2. What `mpg` would you predict for a car with a displacement of 333?
3. What `mpg` would you predict for a car with a displacement of 12 cylinders?
4. What `mpg` would you predict for a car with a displacement of 333 and 12 cylinders?
5. What `mpg` would you predict for a car with a displacement of 333, 12 cylinders, and weighs 4,000 pounds?

14.1.2 `trees`

Open the `trees` dataset in R.

1. What are the variables and what do they mean?
2. Make a plot with Volume on the x axis and Height on the Y and add a best fit line.

3. Use Girth and Height to predict Volume. What would you predict for a tree with a Girth of 10 and a Height of 100 feet?
4. Use Girth and Height to predict Volume. What would you predict for a tree with a Girth of 10 and a Height of 15 meters?
5. What is the maximum circumference of a tree in this dataset?

14.2 More Exercises for you

1. Open the `women` data set. Add a new variable (column) to the `women` dataframe called GPA which is these 15 numbers: 1.5, 3.7, 4.1, 3, 2.5, 3.8, 0.8, 2, 4, 1, 3, 2.5, 3.0, 4.0. You should get something that looks similar to mine.

	height	weight	GPA
FALSE 1	58	115	1.5
FALSE 2	59	117	3.7
FALSE 3	60	120	4.0
FALSE 4	61	123	1.0
FALSE 5	62	126	3.0
FALSE 6	63	129	2.5
FALSE 7	64	132	3.8
FALSE 8	65	135	0.8
FALSE 9	66	139	2.0
FALSE 10	67	142	4.0
FALSE 11	68	146	1.0
FALSE 12	69	150	3.0
FALSE 13	70	154	2.5
FALSE 14	71	159	3.0
FALSE 15	72	164	4.0

2. Use GPA and weight to predict the height of a person who is 155 pounds and has a GPA if 3.33. What is your prediction?
3. Is GPA a significant predictor of height and how do you know?
4. Create a figure showing a best fit line on of height and GPA.
5. Install the `dplyr` package into your Rstudio session.

Chapter 15

Deeper Linear Regression

Let's chat about why understanding linear regression is so important.

While there may always seem to be something new, cool, and shiny in the field of AI/ML, classic statistical methods that leverage machine learning techniques remain powerful and practical for solving many real-world business problems.

Let's look at a very simple model first. For this example, we will need to import the Introduction to Statistical Learning package (ISLR). We will use the “credit” data set that is part of the ISLR package.

```
library(ISLR)
data("Credit")
attach(Credit)

M1 <- lm(Balance ~ Limit + Ethnicity)
```

lm is the function we use to create linear regression models. Now, before we discuss interpreting the results we get from this function, we will discuss the different parts of the model. The “~” symbol is the key to this entire equation. We are telling R to predict whatever is on the left side of the tilde using the variables on the right.

15.1 Interpretation of the Model

Let's run a summary on this model and see what we get.

```
summary(M1)
#>
#> Call:
#> lm(formula = Balance ~ Limit + Ethnicity)
#>
#> Residuals:
#>      Min       1Q   Median       3Q      Max
#> -677.39 -145.75   -8.75  139.56  776.46
#>
#> Coefficients:
#>                Estimate Std. Error t value Pr(>|t|)
#> (Intercept)    -3.078e+02  3.417e+01  -9.007  <2e-16
#> Limit           1.718e-01  5.079e-03  33.831  <2e-16
#> EthnicityAsian   2.835e+01  3.304e+01   0.858    0.391
#> EthnicityCaucasian 1.381e+01  2.878e+01   0.480    0.632
#>
#> (Intercept)      ***
#> Limit            ***
#> EthnicityAsian
#> EthnicityCaucasian
#> ---
#> Signif. codes:
#> 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 234 on 396 degrees of freedom
#> Multiple R-squared:  0.743, Adjusted R-squared:  0.7411
#> F-statistic: 381.6 on 3 and 396 DF, p-value: < 2.2e-16
```

There is a lot of statistical jargon included in our summary that may be unfamiliar to those who have not taken statistics before. That is okay, however, because we are going to breakdown the main statistics we are interested in. Let's start with our variables and their significance in the model.

15.1.1 P-Values

The p-value of our model helps us either prove or disprove the null-hypothesis of our test. In the case of this class, the null-hypothesis is that there is no relationship between the variables we are using to make the predictions and the actual variable we are predicting. In other words, the smaller our p-value the higher the level of significance there is between our variables. When we run a summary of our linear regression model we are given multiple p-values.

First, under coefficients, they are listed for each variable. This can help us optimize our model because we can see what variables are helping make the model more accurate versus those that may be hindering its performance. Also

notice the asterisks next to our p-values. R kindly puts up to three stars next to each variable to help us visually tell if they are significant, essentially more stars means a lower p-value and thus a higher correlation. The second place we see a p-value is at the bottom of our summary. This p-value will give us the overall correlation that exists in our model. As we see in this case, our p-values for this model is $< .00000000000000022$, that is a tiny number and frankly a great p-value. Typically we want our p-value to be .05 or smaller. A p-value of .05 tells us that we have a confidence level of 95%.

15.1.2 Multiple R-Squared

R-squared tells us how well our model explains the variance in our variable. In other words, is the reason for the change in the independent variable actually due to our model's prediction? The higher the r-squared, the more accurate our model is because the better the data fits it. The maximum value r-squared can be is 1.

In our model's case, we have a multiple r-squared of .743, this means our model is approximately 74.3% accurate as this is the amount of variance in the data caused by our dependent variable. Our r-squared could certainly be better. In fact, in the real world you typically are aiming for an r-squared above .9 or .95, which means you would have 90%-95% accuracy.

15.2 Applying the Model to Make Predictions

This type of regression is referred to as linear for a reason. If we were to visualize our model on a quadratic plane, we would see a line of best fit that would travel along through our data. This means we can simplify the model to fit the slope-intercept equation:

$$y = m(x) + b$$

In our case the slope of our line is related to the independent variables. The sum of these slopes will give us the overall slope of our line and the intercept is provided by the equation summary. If we modify this equation to be more applicable to our situation we would get something like this:

$$y = m_1x_1 + m_2x_2 \dots + b$$

Let's look back at our example model from before

```
M1 <- lm(Balance ~ Limit + Ethnicity)
summary(M1)
#>
#> Call:
#> lm(formula = Balance ~ Limit + Ethnicity)
```

```

#>
#> Residuals:
#>      Min       1Q   Median       3Q      Max
#> -677.39 -145.75   -8.75  139.56  776.46
#>
#> Coefficients:
#>                Estimate Std. Error t value Pr(>|t|)
#> (Intercept)    -3.078e+02  3.417e+01  -9.007  <2e-16
#> Limit           1.718e-01  5.079e-03  33.831  <2e-16
#> EthnicityAsian    2.835e+01  3.304e+01   0.858    0.391
#> EthnicityCaucasian 1.381e+01  2.878e+01   0.480    0.632
#>
#> (Intercept)      ***
#> Limit            ***
#> EthnicityAsian
#> EthnicityCaucasian
#> ---
#> Signif. codes:
#> 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 234 on 396 degrees of freedom
#> Multiple R-squared:  0.743, Adjusted R-squared:  0.7411
#> F-statistic: 381.6 on 3 and 396 DF, p-value: < 2.2e-16

```

We see that our limit variable has an estimate of 1.718e-01, this is our slope. When dealing with quantitative variables, we simply multiply our slope by the intended independent variable. So, if we wanted to find the balance of someone with a limit of 400, we would multiply 1.718e-01 by 400.

With the qualitative variables, in this case ethnicity, we multiply the estimate of the TRUE values by 1 and FALSE values by 0, thus cancelling the FALSE values out.

Let's look at an example. If we used our above equation to predict the balance of someone who was Caucasian and has a credit limit of 500, here is the equation we would set up:

```

y = (1.718e-1*500) + (1.381e+01*1) + (2.835e+01*0) + (-3.078e+02)
y
#> [1] -208.09

```

So, according to our model our customer would have a balance of -208.09. This number may seem funny, but keep in mind that our r-squared was not the best for this model making it inaccurate and the ethnicity of the customer was not highly correlated with the balance. Both of these facts may cause our prediction

to be off. If we were actually creating a model that could predict balance, we would want to look at some of the more correlated variables in the data set.

15.3 Review Questions

- 1) Create a linear model predicting using the ISLR data set that predicts a customer's credit limit based on their age, current balance, and the number of cards they have.
- 2) What is the p-value of this model? What does this tell us?
- 3) List the variables in order from most correlated to least. How do you know that they are correlated?
- 4) What is the multiple r-squared of the model? What does this tell us? Is this good or bad?
- 5) What would be the credit limit of a 29 year old with 5 cards and a total balance of 1500?
- 6) Explain what the following piece of code does

```
library(ISLR)
data("Credit")
attach(Credit)
q1 <- lm(Cards ~ Limit + Balance + Education)
```


Chapter 16

Linear Regression Practice

16.1 Your homework is to watch these videos which are posted under the linear regression header on Brightspace and then do the following homework:

16.1.1 Videos to watch:

1. Linear regression women
2. Best fit line women

16.2 Problems

Once you have watched these videos, and you can refer to them as often as you would like, please answer and do the following:

1. Use linear regression to predict the weight of a woman who is 100 inches tall.
2. Use linear regression to predict the height of the woman who weighs 200 pounds.
3. Use linear regression to predict the height of a woman who weighs 5 pounds.
4. Use linear regression to predict the weight of a woman who is 200 inches tall.
5. Plot weight on the X axes and height on the y-axes and create a best fit line on your plot.

6. Plot height on the y-axis and wait on the X axes and create a best fit line on your plot.
7. Add a another column to the women dataframe called GPA which is these 15 numbers: 1.5, 4, 2, 3.7, 4, 1, 3, 2.5, 3.8, 0.8, 2, 4, 1, 3, 2.
8. Use GPA to predict height. Is GPA a significant predictor and how do you know? Draw a best fit line on this relationship.
9. Use GPA to predict a weight. Is GPA a significant predictor and how do you know? Draw a bested line on this relationship, too.
10. Predict the height of a person with a GPA of 4.0.

16.3 Multivariate Regression

I have posted a short video walking you through how to perform multiple linear regression – where you have more than one variable predicting another.

Using the data set mtcars data set:

1. Which variable predicts miles per gallon better gear or qsec? How can you tell?
2. Which two variables out of these four (qsec, vs, am, gear) together best predict miles per gallon?
3. Using only the number of cylinders, displacement, and weight what would mpg you would you predict for a car with a displacement of 400 inches, eight cylinders, and weighing 2000 pounds?
4. Be able to explain in a model which variables are significantly significant.
5. Be able to explain what adjusted R squared means.

Chapter 17

Filters and packages

Filtering data is one of the very basic operation when you work with data. You want to remove a part of the data that is invalid or simply you're not interested in. Or, you want to zero in on a particular part of the data you want to know more about

For example, in the `randu` dataset, how many `y` variables are greater than 0.5? 0.6?

```
length(randu$y[randu$y>0.5])
#> [1] 191
new.randu <- randu$y[randu$y>0.6]
head(new.randu)
#> [1] 0.873416 0.648545 0.826873 0.926590 0.741526 0.846041
length(new.randu)
#> [1] 161
```

In the `randu` dataset, how many `z` variables are greater than 0.9? Less than 0.1? Greater than 0.9 or less than 0.1?

```
length(randu$z[randu$z>0.9])
#> [1] 29
length(randu$z[randu$z<0.1])
#> [1] 37
```

17.1 R packages

From Wikipedia, the free encyclopedia, and fount of all knowledge

R packages are extensions to the R statistical programming language. R packages contain code, data, and documentation in a standardised collection format that can be installed by users of R, typically via a centralised software repository such as CRAN (the Comprehensive R Archive Network).

The large number of packages available for R, and the ease of installing and using them, has been cited as a major factor in driving the widespread adoption of the language in data science.

17.2 You can install the latest released version from CRAN with:

```
install.packages("dplyr")
```

17.3 In RStudio

Installing Packages

- Open RStudio. ...
- In the lower-right pane of RStudio, select the Packages tab and the Install button.
- Type the name of the packages to be installed in the “Packages (separate multiple packages with a space or comma):” box. ...
- Press Install .

17.4 Check this out

dplyr link

<https://dplyr.tidyverse.org/>

Chapter 18

dplyr

18.1 Introduction

For more help **PLEASE** check out Introduction to dplyr introducing the key functionality of the dplyr package.

<https://dplyr.tidyverse.org/articles/dplyr.html>

Your life is about to change. For the better, even.

18.2 A Neat Resource

- RStudio's Data Wrangling Cheat Sheet for dplyr and tidyr

18.3 Single table verbs

dplyr aims to provide a function for each basic verb of data manipulation. These verbs can be organised into three categories based on the component of the dataset that they work with:

Rows:

- **filter()** chooses rows based on column values.
- **slice()** chooses rows based on location.
- **arrange()** changes the order of the rows.

Columns:

- `select()` changes whether or not a column is included.
- `rename()` changes the name of columns. `mutate()` changes the values of columns and creates new columns.
- `relocate()` changes the order of the columns. Groups of rows:
- `summarise()` collapses a group into a single row. It's not that useful until we learn the `group_by()` verb below.

18.4 The pipe

All of the `dplyr` functions take a data frame (or tibble) as the first argument. You can use the pipe to rewrite multiple operations that you can read left-to-right, top-to-bottom (reading the pipe operator as “then”).

What is this: `%>%`?

18.5 Loading dplyr

```
# You should already have done this but you'll need it
install.packages("dplyr")
```

18.6 starwars examples

```
library(dplyr)

starwars %>%
  filter(species == "Droid")
#> # A tibble: 6 x 14
#>   name    height  mass hair_color skin_color eye_color
#>   <chr>    <int> <dbl> <chr>      <chr>      <chr>
#> 1 C-3PO     167    75 <NA>      gold        yellow
#> 2 R2-D2      96    32 <NA>      white, blue red
#> 3 R5-D4      97    32 <NA>      white, red  red
#> 4 IG-88     200   140 none      metal       red
#> 5 R4-P17     96    NA none      silver, red red, blue
#> 6 BB8       NA    NA none      none        black
#> # ... with 8 more variables: birth_year <dbl>, sex <chr>,
#> #   gender <chr>, homeworld <chr>, species <chr>,
#> #   films <list>, vehicles <list>, starships <list>
```

```

starwars %>%
  select(name, ends_with("color"))
#> # A tibble: 87 x 4
#>   name                hair_color skin_color eye_color
#>   <chr>                <chr>      <chr>      <chr>
#> 1 Luke Skywalker      blond        fair        blue
#> 2 C-3PO                <NA>         gold        yellow
#> 3 R2-D2                <NA>         white, blue red
#> 4 Darth Vader         none         white        yellow
#> 5 Leia Organa         brown        light        brown
#> 6 Owen Lars           brown, grey  light        blue
#> 7 Beru Whitesun lars brown        light        blue
#> 8 R5-D4                <NA>         white, red   red
#> 9 Biggs Darklighter   black        light        brown
#> 10 Obi-Wan Kenobi     auburn, white fair        blue-gray
#> # ... with 77 more rows

starwars %>%
  mutate(name, bmi = mass / ((height / 100) ^ 2)) %>%
  select(name:mass, bmi)
#> # A tibble: 87 x 4
#>   name                height mass    bmi
#>   <chr>                <int> <dbl> <dbl>
#> 1 Luke Skywalker      172    77  26.0
#> 2 C-3PO                167    75  26.9
#> 3 R2-D2                96     32  34.7
#> 4 Darth Vader         202   136  33.3
#> 5 Leia Organa         150     49  21.8
#> 6 Owen Lars           178   120  37.9
#> 7 Beru Whitesun lars  165    75  27.5
#> 8 R5-D4                97     32  34.0
#> 9 Biggs Darklighter   183    84  25.1
#> 10 Obi-Wan Kenobi     182    77  23.2
#> # ... with 77 more rows

starwars %>%
  arrange(desc(mass))
#> # A tibble: 87 x 6
#>   name                height mass hair_color skin_color eye_color
#>   <chr>                <int> <dbl> <chr>      <chr>      <chr>
#> 1 Jabba Desil~       175  1358 <NA>        green-tan~ orange
#> 2 Grievous           216   159 none        brown, wh~ green, y~

```

```

#> 3 IG-88      200  140 none      metal      red
#> 4 Darth Vader 202  136 none      white      yellow
#> 5 Tarfful    234  136 brown     brown      blue
#> 6 Owen Lars  178  120 brown, gr~ light      blue
#> 7 Bossk      190  113 none      green      red
#> 8 Chewbacca  228  112 brown     unknown    blue
#> 9 Jek Tono Po~ 180  110 brown     fair       blue
#> 10 Dexter Jett~ 198  102 none      brown      yellow
#> # ... with 77 more rows, and 8 more variables:
#> #   birth_year <dbl>, sex <chr>, gender <chr>,
#> #   homeworld <chr>, species <chr>, films <list>,
#> #   vehicles <list>, starships <list>

starwars %>%
  group_by(species) %>%
  summarise(
    n = n(),
    mass = mean(mass, na.rm = TRUE)
  ) %>%
  filter(
    n > 1,
    mass > 50
  )
#> # A tibble: 8 x 3
#>   species      n mass
#>   <chr>    <int> <dbl>
#> 1 Droid         6  69.8
#> 2 Gungan        3   74
#> 3 Human       35  82.8
#> 4 Kaminoan      2   88
#> 5 Mirialan      2  53.1
#> 6 Twi'lek       2   55
#> 7 Wookiee       2  124
#> 8 Zabrak        2   80

```

Chapter 19

Practice

19.1 `starwars` Exercises

Please use the `starwars` dataset from the `dplyr` package to answer the following questions:

1. How many humans are in this dataset?
2. How many characters are taller than 89 cm?
3. How many characters are taller than 37 inches?
4. How many characters are taller than 37 inches and weigh more than 55 pounds?
5. How many characters are not human or droid?
6. How many characters are not human or droid and are taller than 47 inches?
7. Which species has the most individuals included in this data set?
8. Which species has the tallest individuals on average?
9. What is the tallest individual for each species?
10. Calculate the BMI for each individual and determine which individual has the highest BMI. Use the formula `bmi = mass/((height/100)^2)` to calculate bmi.
11. Which homeworld has the most individuals included in this data set?
12. Which homeworld has the tallest individuals on average?
13. What is the tallest individual for each eye color?

19.2 `iris` Exercises

Please use the `iris` dataset from base R to answer the following questions:

1. How many “virginica” have a petal width of 2.3 or greater in this dataset?

```
#>      n
#> 1 14
```

2. What is the average petal width for each species?

```
#> # A tibble: 3 x 2
#>   Species    avg.petal.width
#>   <fct>         <dbl>
#> 1 setosa         0.246
#> 2 versicolor     1.33
#> 3 virginica       2.03
```

3. How many observations of each species have a petal width of 2.3 or greater in this dataset?

```
#> # A tibble: 1 x 2
#>   Species      n
#>   <fct>    <int>
#> 1 virginica    14
```

4. How many observations of each species have a petal width of 0.5 or greater in this dataset?

```
#> # A tibble: 3 x 2
#>   Species      n
#>   <fct>    <int>
#> 1 setosa         2
#> 2 versicolor    50
#> 3 virginica     50
```

19.3 More practice

Using the `flights` data answer:

1. How many different destinations are in the flights data?
2. Which destination has the most flights?
3. How many origin airports are in the state of sick, and which has the most flights?
4. Which month of the year has the most flights? The least?
5. Which Carrier has the longest average flight?
6. Which carrier has the most flights in the month of October leaving from Newark?

7. What percentage of June flights arrive late?
8. Which carrier has the most late flights in the month of December?
9. This one is hard: for flights leaving JFK, which airline has the greatest percentage of late flights?
10. Think of a question that you could ask about this data. What's the question, and what's the answer?

19.4 Yet more practice

Using the `iris` data answer:

1. How many flowers are species “virginica” and have a `Sepal.Length` greater than 7.5?
2. Create a new variable called `Sepal.Area` with is: `Sepal.Area = Sepal.Width * Sepal.Length`
3. What is the maximum new `Sepal.Area`?
4. Calculate the average `Sepal.Area` for each species and organize these values in descending order.
5. Find the minimum `Sepal.Area` for each species and organize these values in descending order.
6. How many flowers of each species have `Sepal.Area` of less than 15?
7. What is the mean sepal length for flowers with a `Sepal. Area` of less than 15?

Chapter 20

dplyr and nycflights13

```
# load packages
suppressMessages(library(dplyr))
library(nycflights13)

# print the flights dataset from nycflights13
head(flights)
#> # A tibble: 6 x 19
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>      <dbl>
#> 1  2013     1     1     517           515         2
#> 2  2013     1     1     533           529         4
#> 3  2013     1     1     542           540         2
#> 4  2013     1     1     544           545        -1
#> 5  2013     1     1     554           600        -6
#> 6  2013     1     1     554           558        -4
#> # ... with 13 more variables: arr_time <int>,
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>
```

20.1 Choosing columns: select, rename

```
# besides just using select() to pick columns...
flights %>% select(carrier, flight)
```

```
#> # A tibble: 336,776 x 2
#>   carrier flight
#>   <chr>    <int>
#> 1 UA      1545
#> 2 UA      1714
#> 3 AA      1141
#> 4 B6       725
#> 5 DL       461
#> 6 UA     1696
#> 7 B6       507
#> 8 EV     5708
#> 9 B6        79
#> 10 AA      301
#> # ... with 336,766 more rows

# ...you can use the minus sign to hide columns
flights %>% select(-month, -day)
#> # A tibble: 336,776 x 17
#>   year dep_time sched_dep_time dep_delay arr_time
#>   <int>   <int>         <int>     <dbl>   <int>
#> 1  2013     517           515         2     830
#> 2  2013     533           529         4     850
#> 3  2013     542           540         2     923
#> 4  2013     544           545        -1    1004
#> 5  2013     554           600        -6     812
#> 6  2013     554           558        -4     740
#> 7  2013     555           600        -5     913
#> 8  2013     557           600        -3     709
#> 9  2013     557           600        -3     838
#> 10 2013     558           600        -2     753
#> # ... with 336,766 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>

# hide a range of columns
flights %>% select(-(dep_time:arr_delay))

# hide any column with a matching name
flights %>% select(-contains("time"))

# pick columns using a character vector of column names
cols <- c("carrier", "flight", "tailnum")
flights %>% select(one_of(cols))
```

```
#> # A tibble: 336,776 x 3
#>   carrier flight tailnum
#>   <chr>     <int> <chr>
#> 1 UA         1545 N14228
#> 2 UA         1714 N24211
#> 3 AA         1141 N619AA
#> 4 B6          725 N804JB
#> 5 DL          461 N668DN
#> 6 UA         1696 N39463
#> 7 B6          507 N516JB
#> 8 EV         5708 N829AS
#> 9 B6          79  N593JB
#> 10 AA         301 N3ALAA
#> # ... with 336,766 more rows
```

select() can be used to rename columns, though all columns not mentioned are dropped

```
flights %>% select(tail = tailnum)
```

```
#> # A tibble: 336,776 x 1
#>   tail
#>   <chr>
#> 1 N14228
#> 2 N24211
#> 3 N619AA
#> 4 N804JB
#> 5 N668DN
#> 6 N39463
#> 7 N516JB
#> 8 N829AS
#> 9 N593JB
#> 10 N3ALAA
#> # ... with 336,766 more rows
```

rename() does the same thing, except all columns not mentioned are kept

```
flights %>% rename(tail = tailnum)
```

```
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>         <dbl>
#> 1  2013     1     1     517             515           2
#> 2  2013     1     1     533             529           4
#> 3  2013     1     1     542             540           2
#> 4  2013     1     1     544             545          -1
#> 5  2013     1     1     554             600          -6
#> 6  2013     1     1     554             558          -4
#> 7  2013     1     1     555             600          -5
#> 8  2013     1     1     557             600          -3
```

```
#> 9 2013 1 1 557 600 -3
#> 10 2013 1 1 558 600 -2
#> # ... with 336,766 more rows, and 13 more variables:
#> #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
#> #   carrier <chr>, flight <int>, tail <chr>, origin <chr>,
#> #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>
```

20.2 Choosing rows: filter, between, slice, sample_n, top_n, distinct

```
# filter() supports the use of multiple conditions
flights %>% filter(dep_time >= 600, dep_time <= 605)
#> # A tibble: 2,460 x 19
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>      <dbl>
#> 1 2013     1     1     600           600          0
#> 2 2013     1     1     600           600          0
#> 3 2013     1     1     601           600          1
#> 4 2013     1     1     602           610         -8
#> 5 2013     1     1     602           605         -3
#> 6 2013     1     2     600           600          0
#> 7 2013     1     2     600           605         -5
#> 8 2013     1     2     600           600          0
#> 9 2013     1     2     600           600          0
#> 10 2013     1     2     600           600          0
#> # ... with 2,450 more rows, and 13 more variables:
#> #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
#> #   carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>,
#> #   distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>
```

```
# between() is a concise alternative for determining if numeric values fall in a range
flights %>% filter(between(dep_time, 600, 605))
```

```
# side note: is.na() can also be useful when filtering
flights %>% filter(!is.na(dep_time))
```

```
# slice() filters rows by position
flights %>% slice(1000:1005)
```

20.2. CHOOSING ROWS: FILTER, BETWEEN, SLICE, SAMPLE_N, TOP_N, DISTINCT87

```
#> # A tibble: 6 x 19
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>     <dbl>
#> 1  2013     1     2     809             810         -1
#> 2  2013     1     2     810             800          10
#> 3  2013     1     2     811             815          -4
#> 4  2013     1     2     811             815          -4
#> 5  2013     1     2     811             820          -9
#> 6  2013     1     2     815             815           0
#> # ... with 13 more variables: arr_time <int>,
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>

# keep the first three rows within each group
flights %>% group_by(month, day) %>% slice(1:3)
#> # A tibble: 1,095 x 19
#> # Groups:   month, day [365]
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>     <dbl>
#> 1  2013     1     1     517             515           2
#> 2  2013     1     1     533             529           4
#> 3  2013     1     1     542             540           2
#> 4  2013     1     2      42            2359          43
#> 5  2013     1     2     126            2250         156
#> 6  2013     1     2     458             500          -2
#> 7  2013     1     3      32            2359           33
#> 8  2013     1     3      50            2145         185
#> 9  2013     1     3     235            2359         156
#> 10 2013     1     4      25            2359          26
#> # ... with 1,085 more rows, and 13 more variables:
#> #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
#> #   carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>,
#> #   distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>

# sample three rows from each group
flights %>% group_by(month, day) %>% sample_n(3)
#> # A tibble: 1,095 x 19
#> # Groups:   month, day [365]
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>     <dbl>
#> 1  2013     1     1     904             906          -2
```

```

#> 2 2013      1      1    1657      1650      7
#> 3 2013      1      1    1301      1150     71
#> 4 2013      1      2    1652      1652      0
#> 5 2013      1      2     641      645     -4
#> 6 2013      1      2    1720      1700     20
#> 7 2013      1      3    1338      1320     18
#> 8 2013      1      3    1812      1815     -3
#> 9 2013      1      3    1644      1620     24
#> 10 2013     1      4    1342      1330     12
#> # ... with 1,085 more rows, and 13 more variables:
#> #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
#> #   carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>,
#> #   distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>

# keep three rows from each group with the top dep_delay
flights %>% group_by(month, day) %>% top_n(3, dep_delay)
#> # A tibble: 1,108 x 19
#> # Groups:   month, day [365]
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>         <dbl>
#> 1  2013     1     1     848           1835          853
#> 2  2013     1     1    1815           1325          290
#> 3  2013     1     1    2343           1724          379
#> 4  2013     1     2    1412            838          334
#> 5  2013     1     2    1607           1030          337
#> 6  2013     1     2    2131           1512          379
#> 7  2013     1     3    2008           1540          268
#> 8  2013     1     3    2012           1600          252
#> 9  2013     1     3    2056           1605          291
#> 10 2013     1     4    2058           1730          208
#> # ... with 1,098 more rows, and 13 more variables:
#> #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
#> #   carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>,
#> #   distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>

# also sort by dep_delay within each group
flights %>% group_by(month, day) %>% top_n(3, dep_delay) %>% arrange(desc(dep_delay))
#> # A tibble: 1,108 x 19
#> # Groups:   month, day [365]
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>         <dbl>

```


20.2. CHOOSING ROWS: FILTER, BETWEEN, SLICE, SAMPLE_N, TOP_N, DISTINCT89

```
#> 1 2013 1 9 641 900 1301
#> 2 2013 6 15 1432 1935 1137
#> 3 2013 1 10 1121 1635 1126
#> 4 2013 9 20 1139 1845 1014
#> 5 2013 7 22 845 1600 1005
#> 6 2013 4 10 1100 1900 960
#> 7 2013 3 17 2321 810 911
#> 8 2013 6 27 959 1900 899
#> 9 2013 7 22 2257 759 898
#> 10 2013 12 5 756 1700 896
#> # ... with 1,098 more rows, and 13 more variables:
#> #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
#> #   carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>,
#> #   distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>
```

```
# unique rows can be identified using unique() from base R
flights %>% select(origin, dest) %>% unique()
#> # A tibble: 224 x 2
#>   origin dest
#>   <chr> <chr>
#> 1 EWR IAH
#> 2 LGA IAH
#> 3 JFK MIA
#> 4 JFK BQN
#> 5 LGA ATL
#> 6 EWR ORD
#> 7 EWR FLL
#> 8 LGA IAD
#> 9 JFK MCO
#> 10 LGA ORD
#> # ... with 214 more rows
```

```
# dplyr provides an alternative that is more "efficient"
flights %>% select(origin, dest) %>% distinct()
```

```
# side note: when chaining, you don't have to include the parentheses if there are no arguments
flights %>% select(origin, dest) %>% distinct
```

20.3 Excercises

Using the nycflights13 dataset and the dplyr package, answer these questions. Some answers are given in square brackets for you to check your answers.

1. How many flights in Sept were late departing flights? [7815]
2. How many flights in Sept were late departing flights that originated at JFK airport? [2649]
3. How many flights in Sept were late departing flights with an origin of JFK airport and had an destination of anywhere except MIA? [2572]
4. Which carrier had the most flights in this data set? [UA with 58665]
5. Which destination had the most flights in this data set? [ORD with 17283]
6. Which destination had the most flights with departure delays of greater than 60 minutes in this data set? [ORD with 1480]
7. What was the longest arrival delay in this dataset? [1272]
8. Which carrier in September had the most late departing flights? [UA with 1559]
9. Create a variable called total.annoyance which arrival delay plus the departure delay for each flight.
10. Which carrier with more than 10 flights in September had greatest % late departing flights?

20.4 Adding new variables: mutate, transmute, add_rownames

```
# mutate() creates a new variable (and keeps all existing variables)
flights %>% mutate(speed = distance/air_time*60)
#> # A tibble: 336,776 x 20
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>      <dbl>
#> 1  2013     1     1     517           515         2
#> 2  2013     1     1     533           529         4
#> 3  2013     1     1     542           540         2
#> 4  2013     1     1     544           545        -1
#> 5  2013     1     1     554           600        -6
#> 6  2013     1     1     554           558        -4
#> 7  2013     1     1     555           600        -5
#> 8  2013     1     1     557           600        -3
#> 9  2013     1     1     557           600        -3
#> 10 2013     1     1     558           600        -2
#> # ... with 336,766 more rows, and 14 more variables:
#> #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
```

20.4. ADDING NEW VARIABLES: MUTATE, TRANSMUTE, ADD_ROWNUMES91

```
#> #   carrier <chr>, flight <int>, tailnum <chr>,  
#> #   origin <chr>, dest <chr>, air_time <dbl>,  
#> #   distance <dbl>, hour <dbl>, minute <dbl>,  
#> #   time_hour <dtm>, speed <dbl>  
  
# transmute() only keeps the new variables  
flights %>% transmute(speed = distance/air_time*60)  
#> # A tibble: 336,776 x 1  
#>   speed  
#>   <dbl>  
#> 1  370.  
#> 2  374.  
#> 3  408.  
#> 4  517.  
#> 5  394.  
#> 6  288.  
#> 7  404.  
#> 8  259.  
#> 9  405.  
#> 10 319.  
#> # ... with 336,766 more rows
```

```
# example data frame with row names  
mtcars %>% head()  
#>  
#>      mpg cyl  disp  hp  drat    wt  qsec vs am  
#> Mazda RX4      21.0   6  160 110 3.90 2.620 16.46  0  1  
#> Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02  0  1  
#> Datsun 710      22.8   4  108  93 3.85 2.320 18.61  1  1  
#> Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44  1  0  
#> Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0  
#> Valiant        18.1   6  225 105 2.76 3.460 20.22  1  0  
#>  
#>      gear carb  
#> Mazda RX4      4    4  
#> Mazda RX4 Wag  4    4  
#> Datsun 710      4    1  
#> Hornet 4 Drive  3    1  
#> Hornet Sportabout 3    2  
#> Valiant        3    1  
  
# add_rownames() turns row names into an explicit variable  
mtcars %>% add_rownames("model") %>% head()  
#> Warning: `add_rownames()` was deprecated in dplyr 1.0.0.  
#> Please use `tibble::rownames_to_column()` instead.  
#> This warning is displayed once every 8 hours.  
#> Call `lifecycle::last_lifecycle_warnings()` to see where this warning was generated.
```

```
#> # A tibble: 6 x 12
#>   model      mpg   cyl  disp    hp  drat    wt   qsec    vs
#>   <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1 Mazda RX4      21     6  160   110  3.9   2.62  16.5     0
#> 2 Mazda RX4~     21     6  160   110  3.9   2.88  17.0     0
#> 3 Datsun 710    22.8     4  108    93  3.85  2.32  18.6     1
#> 4 Hornet 4 ~    21.4     6  258   110  3.08  3.22  19.4     1
#> 5 Hornet Sp~    18.7     8  360   175  3.15  3.44  17.0     0
#> 6 Valiant      18.1     6  225   105  2.76  3.46  20.2     1
#> # ... with 3 more variables: am <dbl>, gear <dbl>,
#> #   carb <dbl>

# side note: dplyr no longer prints row names (ever) for local data frames
mtcars %>% tbl_df()
#> Warning: `tbl_df()` was deprecated in dplyr 1.0.0.
#> Please use `tibble::as_tibble()` instead.
#> This warning is displayed once every 8 hours.
#> Call `lifecycle::last_lifecycle_warnings()` to see where this warning was generated
#> # A tibble: 32 x 11
#>   mpg   cyl  disp    hp  drat    wt   qsec    vs    am
#>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1  21     6  160   110  3.9   2.62  16.5     0     1
#> 2  21     6  160   110  3.9   2.88  17.0     0     1
#> 3 22.8     4  108    93  3.85  2.32  18.6     1     1
#> 4 21.4     6  258   110  3.08  3.22  19.4     1     0
#> 5 18.7     8  360   175  3.15  3.44  17.0     0     0
#> 6 18.1     6  225   105  2.76  3.46  20.2     1     0
#> 7 14.3     8  360   245  3.21  3.57  15.8     0     0
#> 8 24.4     4  147.    62  3.69  3.19  20.0     1     0
#> 9 22.8     4  141.    95  3.92  3.15  22.9     1     0
#> 10 19.2     6  168.   123  3.92  3.44  18.3     1     0
#> # ... with 22 more rows, and 2 more variables: gear <dbl>,
#> #   carb <dbl>
```

20.5 Grouping and counting: summarise, tally, count, group_size, n_groups, ungroup

```
# summarise() can be used to count the number of rows in each group
flights %>% group_by(month) %>% summarise(cnt = n())
#> # A tibble: 12 x 2
#>   month    cnt
#>   <int> <int>
```

20.5. GROUPING AND COUNTING: SUMMARISE, TALLY, COUNT, GROUP_SIZE, N_GROUPS, UNGROUP

```
#> 1      1 27004
#> 2      2 24951
#> 3      3 28834
#> 4      4 28330
#> 5      5 28796
#> 6      6 28243
#> 7      7 29425
#> 8      8 29327
#> 9      9 27574
#> 10     10 28889
#> 11     11 27268
#> 12     12 28135
```

```
# tally() and count() can do this more concisely
flights %>% group_by(month) %>% tally()
flights %>% count(month)
```

```
# you can sort by the count
flights %>% group_by(month) %>% summarise(cnt = n()) %>% arrange(desc(cnt))
#> # A tibble: 12 x 2
#>   month   cnt
#>   <int> <int>
#> 1     7 29425
#> 2     8 29327
#> 3    10 28889
#> 4     3 28834
#> 5     5 28796
#> 6     4 28330
#> 7     6 28243
#> 8    12 28135
#> 9     9 27574
#> 10    11 27268
#> 11     1 27004
#> 12     2 24951
```

```
# tally() and count() have a sort parameter for this purpose
flights %>% group_by(month) %>% tally(sort=TRUE)
flights %>% count(month, sort=TRUE)
```

```
# you can sum over a specific variable instead of simply counting rows
flights %>% group_by(month) %>% summarise(dist = sum(distance))
#> # A tibble: 12 x 2
#>   month   dist
#>   <int>   <dbl>
```

```
#> 1      1 27188805
#> 2      2 24975509
#> 3      3 29179636
#> 4      4 29427294
#> 5      5 29974128
#> 6      6 29856388
#> 7      7 31149199
#> 8      8 31149334
#> 9      9 28711426
#> 10     10 30012086
#> 11     11 28639718
#> 12     12 29954084
```

```
# tally() and count() have a wt parameter for this purpose
flights %>% group_by(month) %>% tally(wt = distance)
flights %>% count(month, wt = distance)
```

```
# group_size() returns the counts as a vector
flights %>% group_by(month) %>% group_size()
#> [1] 27004 24951 28834 28330 28796 28243 29425 29327 27574
#> [10] 23889 27268 28135
```

```
# n_groups() simply reports the number of groups
flights %>% group_by(month) %>% n_groups()
#> [1] 12
```

```
# group by two variables, summarise, arrange (output is possibly confusing)
flights %>% group_by(month, day) %>% summarise(cnt = n()) %>% arrange(desc(cnt)) %>% print()
#> `summarise()` has grouped output by 'month'. You can
#> override using the `.groups` argument.
#> # A tibble: 365 x 3
#> # Groups:   month [12]
#>   month   day   cnt
#>   <int> <int> <int>
#> 1     11    27  1014
#> 2      7    11  1006
#> 3      7     8  1004
#> 4      7    10  1004
#> 5     12     2  1004
#> 6      7    18  1003
#> 7      7    25  1003
#> 8      7    12  1002
#> 9      7     9  1001
#> 10     7    17  1001
```

20.5. GROUPING AND COUNTING: SUMMARISE, TALLY, COUNT, GROUP_SIZE, N_GROUPS, UNGROUP

```
#> 11      7     31 1001
#> 12      8      7 1001
#> 13      8      8 1001
#> 14      8     12 1001
#> 15      7     22 1000
#> 16      7     24 1000
#> 17      8      1 1000
#> 18      8      5 1000
#> 19      8     15 1000
#> 20     11     21 1000
#> 21      7     15 999
#> 22      7     19 999
#> 23      7     26 999
#> 24      7     29 999
#> 25      8      2 999
#> 26      8      9 999
#> 27     11     22 999
#> 28      8     16 998
#> 29      7     23 997
#> 30      7     30 997
#> 31      8     14 997
#> 32      7     16 996
#> 33      8      6 996
#> 34      8     19 996
#> 35      9     13 996
#> 36      9     26 996
#> 37      9     27 996
#> 38      4     15 995
#> 39      6     20 995
#> 40      6     26 995
#> # ... with 325 more rows

# ungroup() before arranging to arrange across all groups
flights %>% group_by(month, day) %>% summarise(cnt = n()) %>% ungroup() %>% arrange(desc(cnt))
#> `summarise()` has grouped output by 'month'. You can
#> override using the `.groups` argument.
#> # A tibble: 365 x 3
#>   month   day   cnt
#>   <int> <int> <int>
#> 1     11    27  1014
#> 2      7     11  1006
#> 3      7      8  1004
#> 4      7     10  1004
#> 5     12      2  1004
#> 6      7     18  1003
```

```
#> 7      7      25 1003
#> 8      7      12 1002
#> 9      7       9 1001
#> 10     7      17 1001
#> # ... with 355 more rows
```

20.6 Creating data frames: `data_frame`

`data_frame()` is a better way than `data.frame()` for creating data frames. Benefits of `data_frame()`:

- You can use previously defined columns to compute new columns.
- It never coerces column types.
- It never munges column names.
- It never adds row names.
- It only recycles length 1 input.
- It returns a local data frame (a `tbl_df`).

```
# data_frame() example
data_frame(a = 1:6, b = a*2, c = 'string', 'd+e' = 1) %>% glimpse()
#> Warning: `data_frame()` was deprecated in tibble 1.1.0.
#> Please use `tibble()` instead.
#> This warning is displayed once every 8 hours.
#> Call `lifecycle::last_lifecycle_warnings()` to see where this warning was generated
#> Rows: 6
#> Columns: 4
#> $ a      <int> 1, 2, 3, 4, 5, 6
#> $ b      <dbl> 2, 4, 6, 8, 10, 12
#> $ c      <chr> "string", "string", "string", "string", "string", "string"
#> $ `d+e` <dbl> 1, 1, 1, 1, 1, 1

# data.frame() example
data.frame(a = 1:6, c = 'string', 'd+e' = 1) %>% glimpse()
#> Rows: 6
#> Columns: 3
#> $ a      <int> 1, 2, 3, 4, 5, 6
#> $ c      <chr> "string", "string", "string", "string", "string", "string"
#> $ d.e <dbl> 1, 1, 1, 1, 1, 1
```


Chapter 21

Joining tables

21.1 Joining (merging) tables: left_join, right_join, inner_join, full_join, semi_join, anti_join

```
# create two simple data frames
(a <- data_frame(color = c("green","yellow","red"), num = 1:3))
#> # A tibble: 3 x 2
#>   color    num
#>   <chr>  <int>
#> 1 green     1
#> 2 yellow    2
#> 3 red       3
(b <- data_frame(color = c("green","yellow","pink"), size = c("S","M","L")))
#> # A tibble: 3 x 2
#>   color size
#>   <chr> <chr>
#> 1 green S
#> 2 yellow M
#> 3 pink  L

# only include observations found in both "a" and "b" (automatically joins on variables that appear in both)
inner_join(a, b)
#> Joining, by = "color"
#> # A tibble: 2 x 3
#>   color    num size
#>   <chr>  <int> <chr>
#> 1 green     1 S
```

```

#> 2 yellow      2 M

# include observations found in either "a" or "b"
full_join(a, b)
#> Joining, by = "color"
#> # A tibble: 4 x 3
#>   color      num size
#>   <chr>   <int> <chr>
#> 1 green         1 S
#> 2 yellow        2 M
#> 3 red           3 <NA>
#> 4 pink          NA L

# include all observations found in "a"
left_join(a, b)
#> Joining, by = "color"
#> # A tibble: 3 x 3
#>   color      num size
#>   <chr>   <int> <chr>
#> 1 green         1 S
#> 2 yellow        2 M
#> 3 red           3 <NA>

# include all observations found in "b"
right_join(a, b)
#> Joining, by = "color"
#> # A tibble: 3 x 3
#>   color      num size
#>   <chr>   <int> <chr>
#> 1 green         1 S
#> 2 yellow        2 M
#> 3 pink          NA L

# right_join(a, b) is identical to left_join(b, a) except for column ordering
left_join(b, a)
#> Joining, by = "color"
#> # A tibble: 3 x 3
#>   color size      num
#>   <chr> <chr>   <int>
#> 1 green S           1
#> 2 yellow M          2
#> 3 pink  L           NA

# filter "a" to only show observations that match "b"
semi_join(a, b)

```

```
#> Joining, by = "color"
#> # A tibble: 2 x 2
#>   color    num
#>   <chr> <int>
#> 1 green     1
#> 2 yellow    2

# filter "a" to only show observations that don't match "b"
anti_join(a, b)
#> Joining, by = "color"
#> # A tibble: 1 x 2
#>   color    num
#>   <chr> <int>
#> 1 red      3

# sometimes matching variables don't have identical names
b <- b %>% rename(col = color)

# specify that the join should occur by matching "color" in "a" with "col" in "b"
inner_join(a, b, by=c("color" = "col"))
#> # A tibble: 2 x 3
#>   color    num size
#>   <chr> <int> <chr>
#> 1 green     1 S
#> 2 yellow    2 M
```

21.2 Viewing more output: print, View

```
# specify that you want to see more rows
flights %>% print(n = 15)
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>         <dbl>
#> 1  2013     1     1     517             515           2
#> 2  2013     1     1     533             529           4
#> 3  2013     1     1     542             540           2
#> 4  2013     1     1     544             545          -1
#> 5  2013     1     1     554             600          -6
#> 6  2013     1     1     554             558          -4
#> 7  2013     1     1     555             600          -5
#> 8  2013     1     1     557             600          -3
#> 9  2013     1     1     557             600          -3
```

```
#> 10 2013      1      1      558      600      -2
#> 11 2013      1      1      558      600      -2
#> 12 2013      1      1      558      600      -2
#> 13 2013      1      1      558      600      -2
#> 14 2013      1      1      558      600      -2
#> 15 2013      1      1      559      600      -1
#> # ... with 336,761 more rows, and 13 more variables:
#> #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
#> #   carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>,
#> #   distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>
```

```
# specify that you want to see ALL rows (don't run this!)
flights %>% print(n = Inf)
```

```
# specify that you want to see all columns
flights %>% print(width = Inf)
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay
#>   <int> <int> <int>   <int>         <int>      <dbl>
#> 1  2013     1     1     517           515         2
#> 2  2013     1     1     533           529         4
#> 3  2013     1     1     542           540         2
#> 4  2013     1     1     544           545        -1
#> 5  2013     1     1     554           600        -6
#> 6  2013     1     1     554           558        -4
#> 7  2013     1     1     555           600        -5
#> 8  2013     1     1     557           600        -3
#> 9  2013     1     1     557           600        -3
#> 10 2013     1     1     558           600        -2
#>   arr_time sched_arr_time arr_delay carrier flight tailnum
#>   <int>         <int>      <dbl> <chr>   <int> <chr>
#> 1     830           819         11 UA      1545 N14228
#> 2     850           830         20 UA      1714 N24211
#> 3     923           850         33 AA      1141 N619AA
#> 4    1004          1022        -18 B6       725 N804JB
#> 5     812           837        -25 DL       461 N668DN
#> 6     740           728         12 UA      1696 N39463
#> 7     913           854         19 B6       507 N516JB
#> 8     709           723        -14 EV      5708 N829AS
#> 9     838           846         -8 B6        79 N593JB
#> 10    753           745          8 AA       301 N3ALAA
#>   origin dest  air_time distance  hour minute
#>   <chr> <chr>    <dbl>    <dbl> <dbl> <dbl>
```

```

#> 1 EWR IAH 227 1400 5 15
#> 2 LGA IAH 227 1416 5 29
#> 3 JFK MIA 160 1089 5 40
#> 4 JFK BQN 183 1576 5 45
#> 5 LGA ATL 116 762 6 0
#> 6 EWR ORD 150 719 5 58
#> 7 EWR FLL 158 1065 6 0
#> 8 LGA IAD 53 229 6 0
#> 9 JFK MCO 140 944 6 0
#> 10 LGA ORD 138 733 6 0
#> time_hour
#> <dtm>
#> 1 2013-01-01 05:00:00
#> 2 2013-01-01 05:00:00
#> 3 2013-01-01 05:00:00
#> 4 2013-01-01 05:00:00
#> 5 2013-01-01 06:00:00
#> 6 2013-01-01 05:00:00
#> 7 2013-01-01 06:00:00
#> 8 2013-01-01 06:00:00
#> 9 2013-01-01 06:00:00
#> 10 2013-01-01 06:00:00
#> # ... with 336,766 more rows

```

```

# show up to 1000 rows and all columns
flights %>% View()

# set option to see all columns and fewer rows
options(dplyr.width = Inf, dplyr.print_min = 6)

# reset options (or just close R)
options(dplyr.width = NULL, dplyr.print_min = 10)

```


Chapter 22

More DPLYR Excerics

Using the `nycflights13` dataset and the `dplyr` package, answer these questions. Some answers are given in square brackets for you to check your answers.

1. How many flights in Sept were late departing flights? [7815]
2. How many flights in Sept were late departing flights that originated at JFK airport? [2649]
3. How many flights in Sept were late departing flights with an origin of JFK airport and had an destination of anywhere except MIA? [2572]
4. Which carrier had the most flights in this data set? [UA with 58665]
5. Which destination had the most flights in this data set? [ORD with 17283]
6. Which destination had the most flights with departure delays of greater than 60 minutes in this data set? [ORD with 1480]
7. What was the longest arrival delay in this dataset? [1272]
8. Which carrier in September had the most late departing flights? [UA with 1559]
9. Create a variable called `total.annoyance` which arrival delay plus the departure delay for each flight.
10. Which carrier with more than 10 flights in September had greatest % late departing flights?

Chapter 23

NFL Data

Please make sure you can load and use the NFL data. You can get the data from:

<https://www.kaggle.com/datasets/maxhorowitz/nflplaybyplay2009to2016?resource=download>

I have also put the data on Brightspace if that helps you.

For Friday, please try to answer these questions:

- Which game had the most plays?
- How many plays?
- Which teams played in this game?
- How many plays if you remove the “No Play” plays?
- Highest scoring game (total points including both teams)?
- How many plays were in this game?

23.1 Additional practice

1. Create a new small data set called `small` that only uses `game_id`'s that are less than 2010000000 and the variables `game_id`, `home_team`, `total_home_score`, `away_team`, `total_away_score`, and `play_type`. How many rows do you get in this data set? I get 41,880.
2. How many plays are in the game 2009111507? I get 167.
3. For `small` how many plays are kick offs? I get 2327.

4. For game 2009111507 how many plays are kick offs? I get 10.
5. For `small` how many plays are qb kneels or qb spikes? I get 425.
6. For `small` how many what is the percentage of run to pass plays excluding all other types of plays? I get 43.15482% running.
7. In `small`, which game has the most passing plays? I get game 2009122006 with 101 passing plays.


Chapter 24

Palmer Penguins Review

24.1 Use the `penguins` dataset from the `palmerpenguins` package for the following:

1. Extract a data frame that excludes the Adelie penguins. Call this data frame `q1`.
2. Extract a data frame that excludes the Adelie penguins and retains those with bill lengths between 40 and 50 mm. Sort it in decreasing order by bill length. Call this data frame `q2`.
3. Calculate the mean bill length and bill depth for each of the three species of penguins in one step (not three different calculations). Do these statistics line up with what you see in the penguins plot below?
4. Consider a new metric called `bill_size` that's the sum of the length and depth. What is the average bill size and its standard deviation among each species, broken out among each of the islands? (that is, nine pairs of statistics) Sort your resulting data frame in decreasing order by average bill size.
5. Answer in words (after you have the numerical answers):
 - Which island has the biggest penguins?
 - Which penguins species are the biggest?
 - Does the size of a penguin depend on the island it lives on?
 - Does the size of each species of penguin depend on the island it lives on?
6. Create a linear model using bill length and bill depth to predict body mass. What would you predict for a penguin with bill length of 50 mm and a bill depth of 23 mm? What about a Gentoo with the same bill dimensions?

7. Create the plot below (or very similar). *FYI, in case these questions are in black and white, each penguin species should be a different color in the plot.*



Chapter 25

Diamonds Practice

25.1 Use the `diamonds` data set found in the `ggplot2` package to answer these questions. You may want to review the data before starting using the help options.

1. How many diamonds in this set have a cut considered ideal?

```
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1 21551
```

2. How many diamonds in this set have a cut considered ideal and have a color of E?

```
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1  3903
```

3. 2. How many diamonds in this set have a cut considered ideal and have a color of E or D?

```
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1  6737
```

4. Organize the average price by cut in descending order.

```
#> # A tibble: 5 x 2
#>   cut      mean
#>   <ord>   <dbl>
#> 1 Premium 4584.
#> 2 Fair    4359.
#> 3 Very Good 3982.
#> 4 Good    3929.
#> 5 Ideal   3458.
```

5. Determine the average price and standard deviation for `ideal` cut diamonds.


```
#> # A tibble: 1 x 2
#>   mean std_dev
#>   <dbl> <dbl>
#> 1 3458.  3808.
```

6. Organize the average price by cut and color in descending order.


```
#> # A tibble: 35 x 3
#> # Groups:   cut [5]
#>   cut      color price
#>   <ord>   <ord> <dbl>
#> 1 Premium J      6295.
#> 2 Premium I      5946.
#> 3 Very Good I      5256.
#> 4 Premium H      5217.
#> 5 Fair    H      5136.
#> 6 Very Good J      5104.
#> 7 Good    I      5079.
#> 8 Fair    J      4976.
#> 9 Ideal   J      4918.
#> 10 Fair   I      4685.
#> # ... with 25 more rows
```

7. Use `ggplot2` to make a plot similar to this:

25.1. USE THE DIAMONDS DATA SET FOUND IN THE GGPLYOT2 PACKAGE TO ANSWER THESE QUESTIONS. Y



8. 7. Use ggplot2 to make a plot similar to this:



9. Predict price of a 2.3 carat diamond with a table of 70 and a depth of 55.

```
#> [1] 15447.52
```

10. In descending order, how many diamonds are there of each clarity?

```
#> # A tibble: 8 x 2
#>   clarity      n
#>   <ord>    <int>
#> 1 SI1      13065
#> 2 VS2      12258
#> 3 SI2       9194
#> 4 VS1       8171
#> 5 VVS2       5066
#> 6 VVS1       3655
#> 7 IF        1790
#> 8 I1         741
```